


CS194/294-129: Designing, Visualizing and Understanding Deep Neural Networks

John Canny

Spring 2018

Lecture 14: Translation

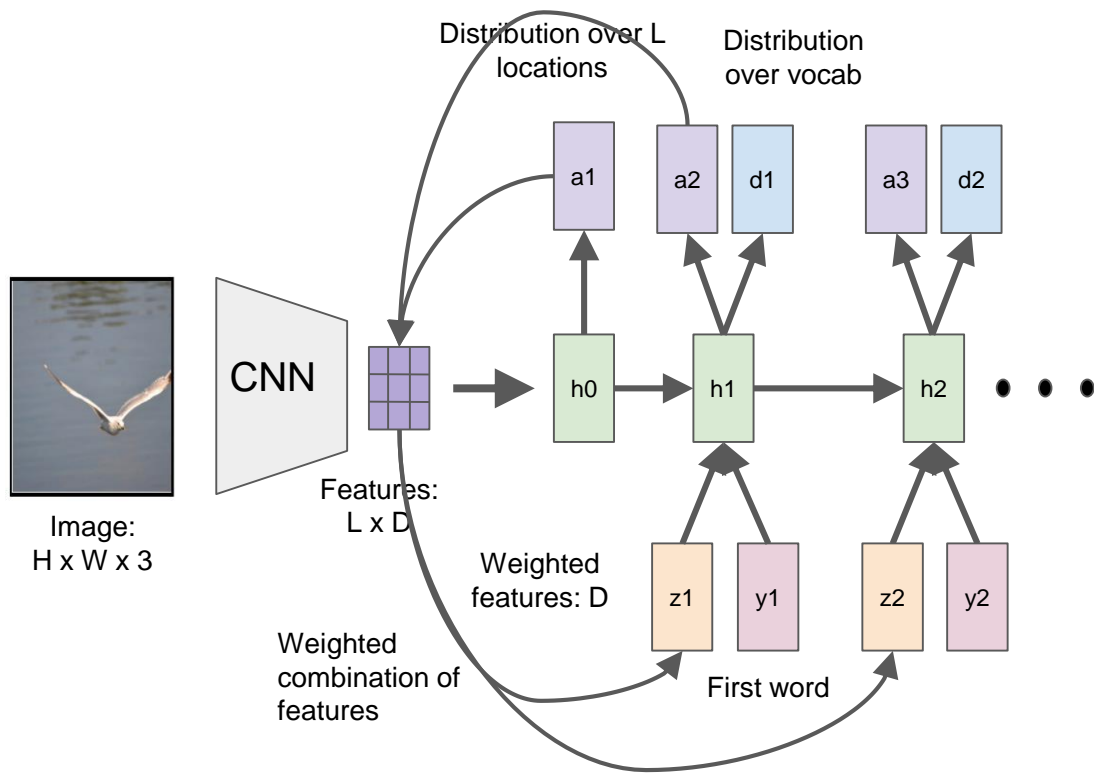
Last Time: Soft vs Hard Attention



Hard attention: Attend to a single input location, can't use gradient descent, Need reinforcement learning.

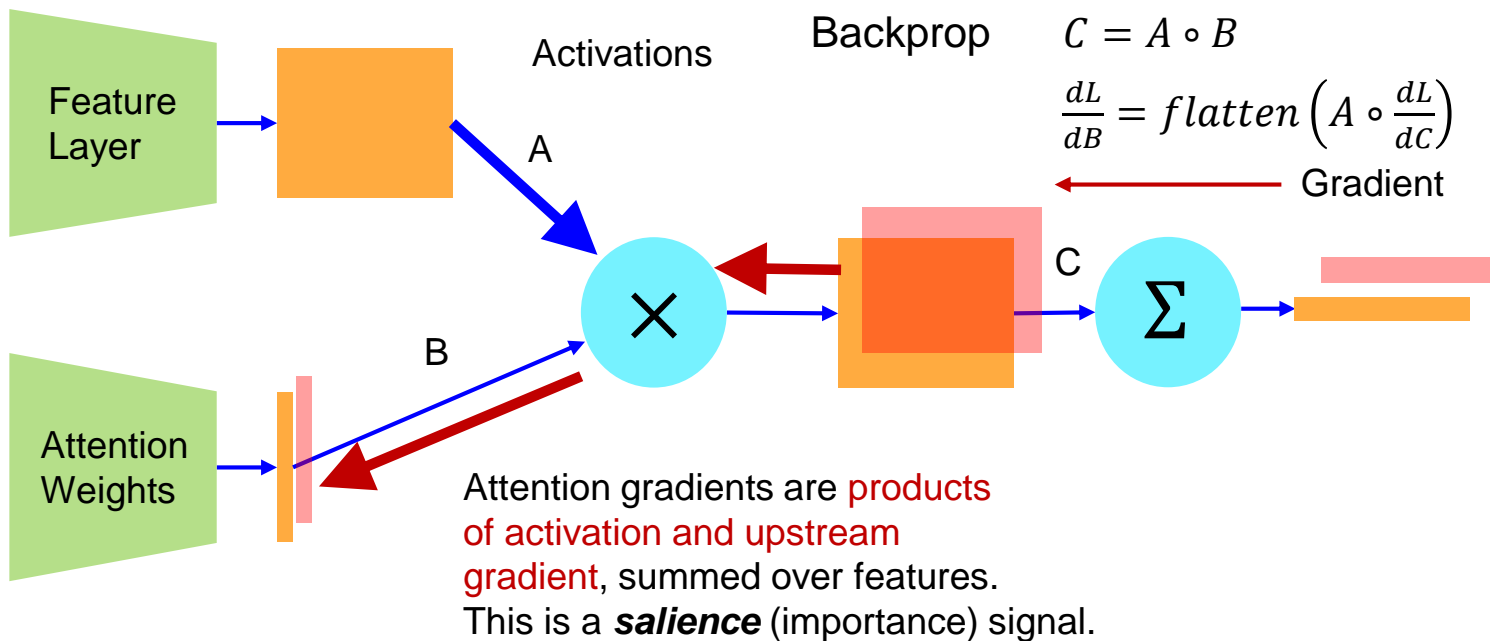
Soft attention: Compute a weighted combination (attention) over some inputs using an attention network. Can use backpropagation to train end-to-end.

Last Time: Recurrent Attention for Captioning



Last Time: Attention Mechanics: Saliency

During training, the attention layer receives gradients which are the **product of the upstream gradient and the feature layer activations** (saliency).

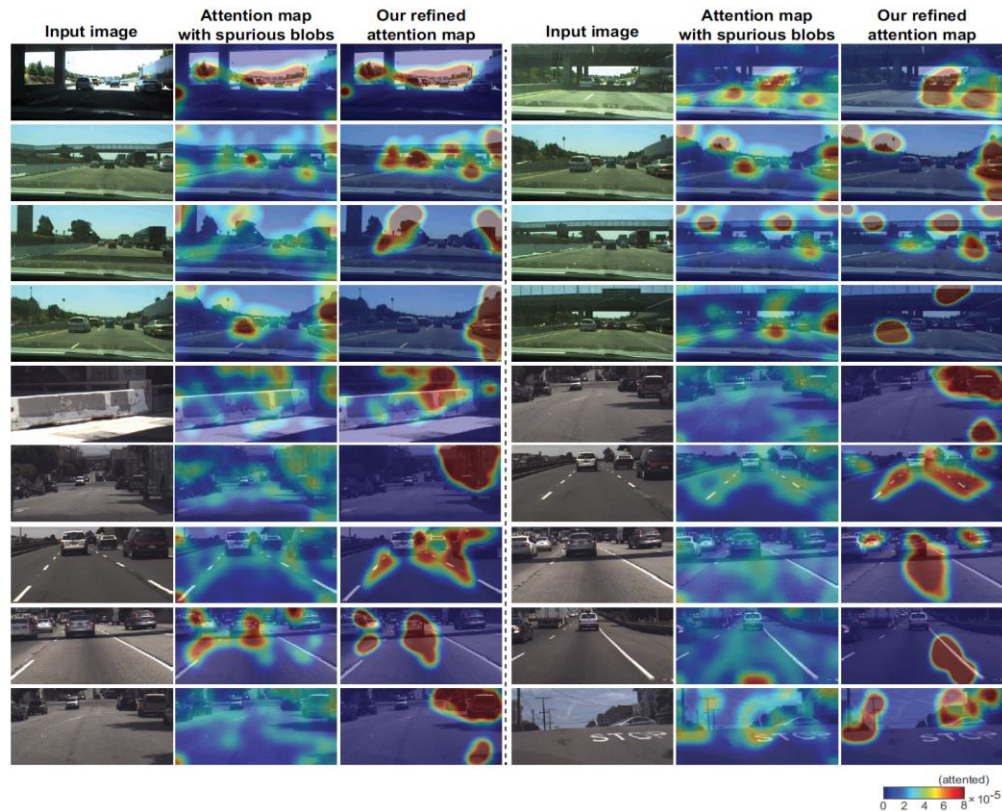


Last Time: Attention and Interpretability

Attention models learn to predict salient (important) inputs.

Attention visualizations help users understand the causes of the network's behavior.

Not every attended region is actually important, but post-processing can remove regions that aren't.




This Time: Translation

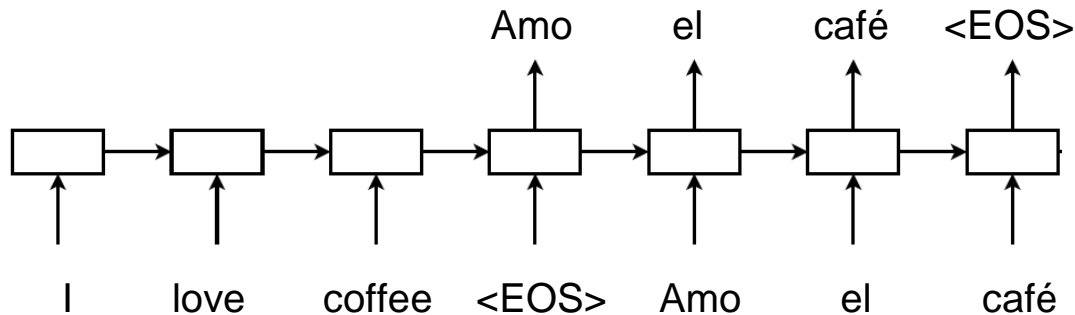
- Sequence-to-sequence translation
- Adding Attention
- Parsing as translation
- Attention only models
- English-to-English translation ?!

Sequence-To-Sequence RNNs

An input sequence is fed to the left array, output sentence to the right array for training:




For translation:



Sequence-To-Sequence RNNs

Generation:



Keep an n-best list of partial sentences, along with their partial softmax scores.

Bleu scores for Translation

The goal of bleu scores is to compare machine translations against human-generated translations, allowing for variation.

Consider these translations for a Chinese sentence:

Candidate 1: It is a guide to action which ensures that the military always obeys the commands of the party.

Candidate 2: It is to insure the troops forever hearing the activity guidebook that party direct.

Bleu Scores for Translation

Candidate 1: It is a guide to action which ensures that the military always obeys the commands of the party.

Candidate 2: It is to insure the troops forever hearing the activity guidebook that party direct.

Reference 1: It is a guide to action that ensures that the military will forever heed Party commands.

Reference 2: It is the guiding principle which guarantees the military forces always being under the command of the Party.

Reference 3: It is the practical guide for the army always to heed the directions of the party.

Bleu Scores for Translation

Candidate 1: It is a guide to action which ensures that the military always obeys the commands of the party.

Candidate 2: It is to insure the troops forever hearing the activity guidebook that party direct.

Reference 1: It is a guide to action that ensures that the military will forever heed Party commands.

Reference 2: It is the guiding principle which guarantees the military forces always being under the command of the Party.

Reference 3: It is the practical guide for the army always to heed the directions of the party.

Bleu Scores for Translation

Candidate 1: It is a guide to action which ensures that the military always obeys the commands of the party.

Candidate 2: It is to insure the troops forever hearing the activity guidebook that party direct.

Reference 1: It is a guide to action that ensures that the military will forever heed Party commands.

Reference 2: It is the guiding principle which guarantees the military forces always being under the command of the Party.

Reference 3: It is the practical guide for the army always to heed the directions of the party

Bleu Scores for Translation

Unigram precision:

$$\frac{\text{correct unigrams occurring in reference sentence}}{\text{unigrams occurring in test sentence}}$$

Modified unigram precision: clip counts by maximum occurrence in any reference sentence:

Candidate: the the the the the the the.

Reference 1: The cat is on the mat.

Reference 2: There is a cat on the mat.

Modified precision is 2/7.

Bleu Scores for Translation

Candidate 1: It is a guide to action which ensures that the military always obeys the commands of the party. **unigram precision 17/18**

Candidate 2: It is to insure the troops forever hearing the activity guidebook that party direct. **unigram precision 8/14**

Reference 1: It is a guide to action that ensures that the military will forever heed Party commands.

Reference 2: It is the guiding principle which guarantees the military forces always being under the command of the Party.

Reference 3: It is the practical guide for the army always to heed the directions of the party.

Bleu Scores for Translation

N-gram precision is defined similarly:

$$\frac{\text{correct ngrams occuring in reference sentence}}{\text{ngrams occuring in test sentence}}$$

Modified ngram precision: clip counts by maximum occurrence in any reference sentence.

Unigram scores tend to capture *adequacy*

Ngram scores tend to capture *fluency*

Bleu Scores for Translation

Candidate 1: It is a guide to action which ensures that the military always obeys the commands of the party. **bigram precision 10/17**

Candidate 2: It is to insure the troops forever hearing the activity guidebook that party direct. **bigram precision 1/13**

Reference 1: It is a guide to action that ensures that the military will forever heed Party commands.

Reference 2: It is the guiding principle which guarantees the military forces always being under the command of the Party.

Reference 3: It is the practical guide for the army always to heed the directions of the party.

Bleu Scores for Translation

How to combine scores for different n-grams?

Averaging sounds good, but precisions are very different for different n (unigrams have much higher scores).

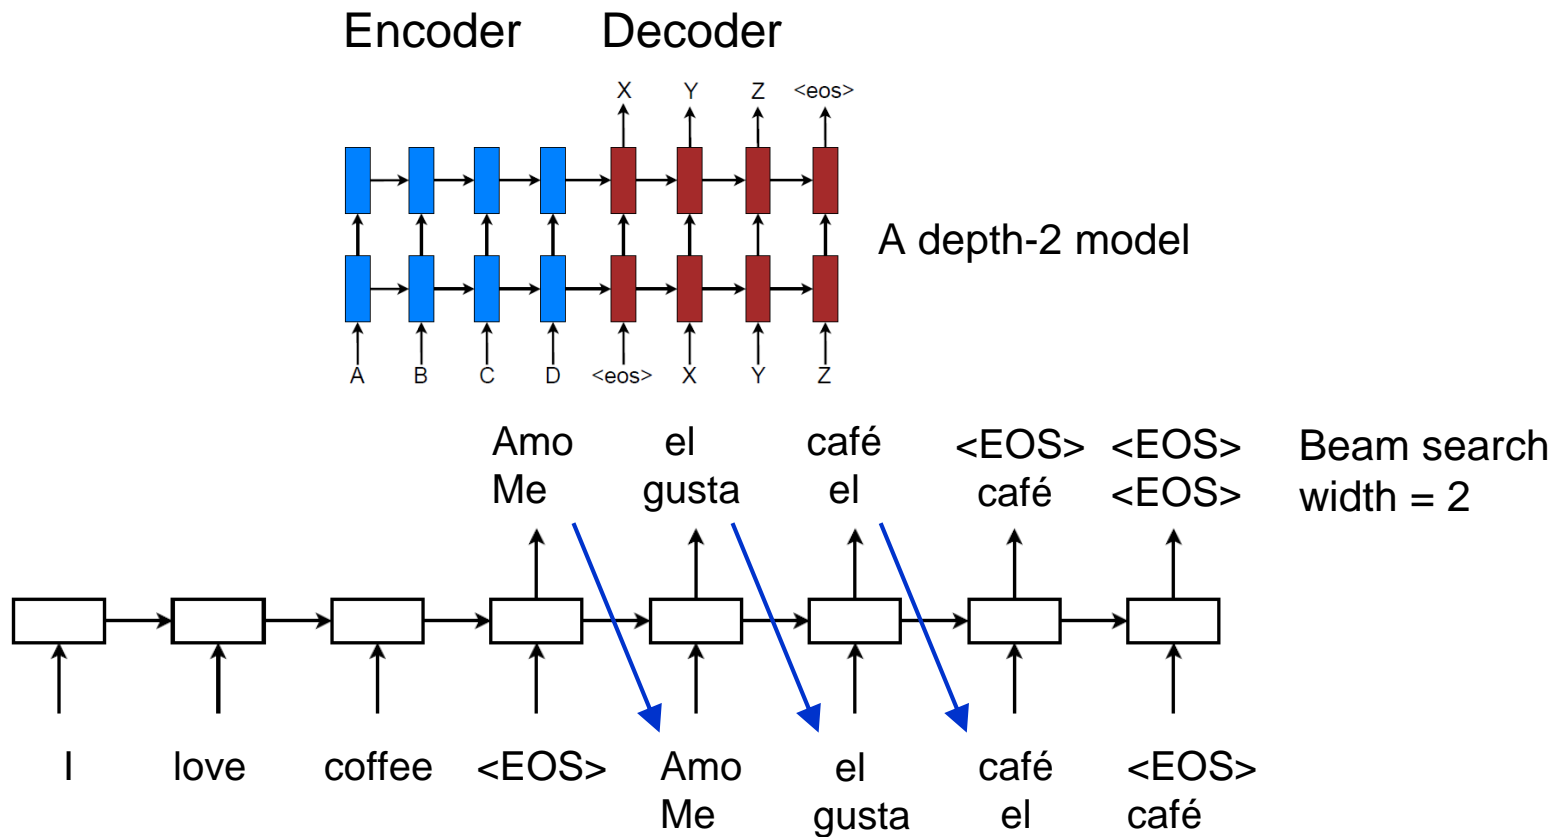
BLEU Score: Take a weighted geometric mean of the logs of n-gram precisions up to some length (usually 4). Add a penalty for too-short predictions.

$$\text{BLEU} = \text{BP} \cdot \exp \left(\sum_{n=1}^N w_n \log p_n \right)$$

$$\text{BP} = \begin{cases} 1 & \text{if } c > r \\ e^{(1-r/c)} & \text{if } c \leq r \end{cases}$$

Candidate length c shorter than reference r translation

Sequence-To-Sequence Model Translation



Sequence-To-Sequence Model Translation

Raw scores for French-English Translation, depth = 4

Method	test BLEU score (ntst14)
Bahdanau et al. [2]	28.45
Baseline System [29]	33.30
Single forward LSTM, beam size 12	26.17
Single reversed LSTM, beam size 12	30.59
Ensemble of 5 reversed LSTMs, beam size 1	33.00
Ensemble of 2 reversed LSTMs, beam size 12	33.27
Ensemble of 5 reversed LSTMs, beam size 2	34.50
Ensemble of 5 reversed LSTMs, beam size 12	34.81

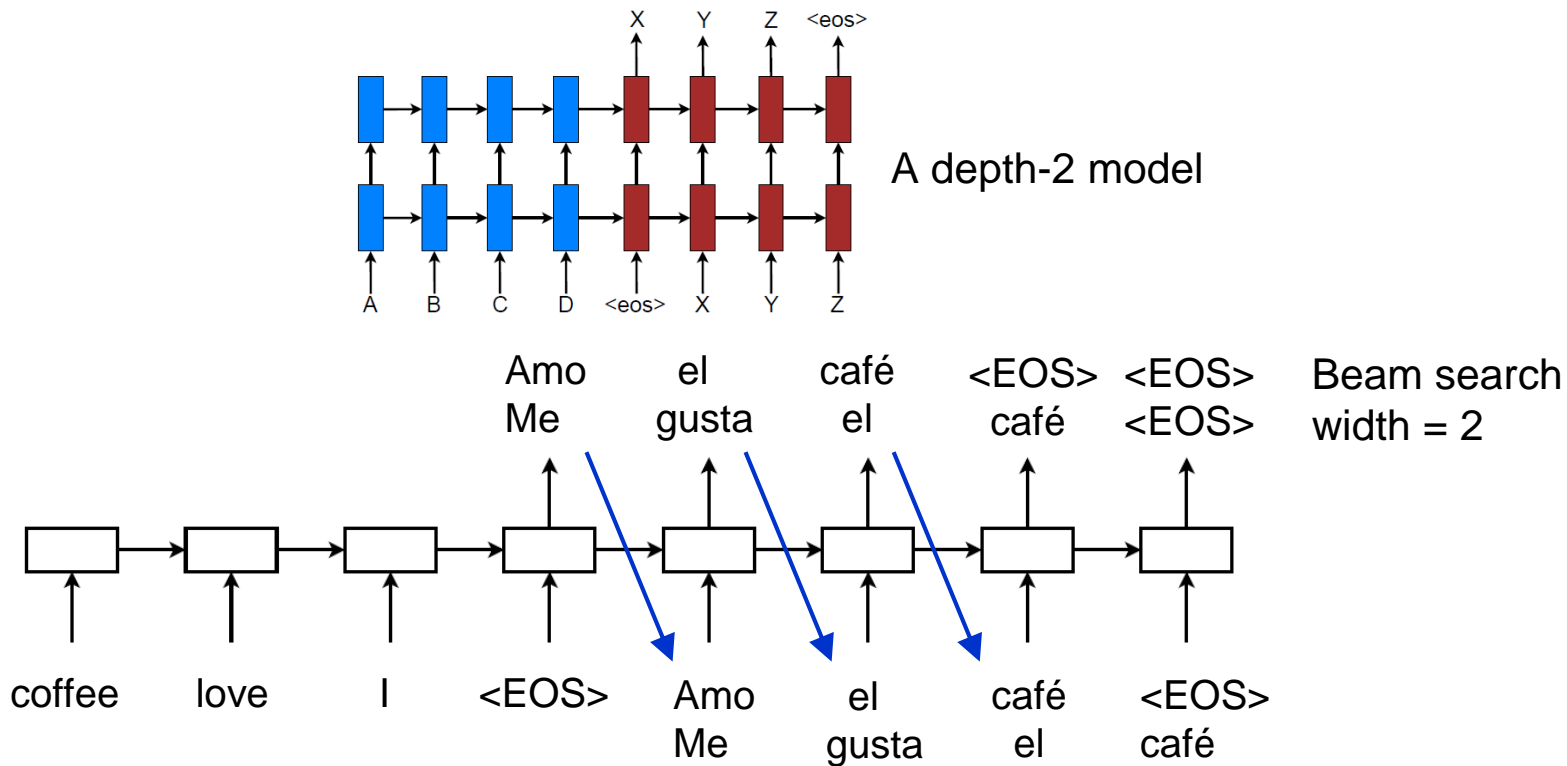
Reversed = reverse the order of the input sentence.

Intuition: the first part of the sentence is the most important, and reversal eases the long-term dependencies from output to input sentence.

From Sutskeyver et al. "Sequence to Sequence Learning with Neural Networks" 2014.

Sequence-To-Sequence Model Translation

Input sequence reversal



Sequence-To-Sequence Model Translation

Raw scores for French-English Translation, depth = 4

Method	test BLEU score (ntst14)
Bahdanau et al. [2]	28.45
Baseline System [29]	33.30
Single forward LSTM, beam size 12	26.17
Single reversed LSTM, beam size 12	30.59
Ensemble of 5 reversed LSTMs, beam size 1	33.00
Ensemble of 2 reversed LSTMs, beam size 12	33.27
Ensemble of 5 reversed LSTMs, beam size 2	34.50
Ensemble of 5 reversed LSTMs, beam size 12	34.81


Beam sizes are tiny!!

The model produces state-of-the-art translations with almost no search.

From Sutskeyver et al. "Sequence to Sequence Learning with Neural Networks" 2014.

Sequence-To-Sequence Criticisms

All the information from the source sentence has to pass through the bottleneck at the last unit(s) of the encoder.




Sentence length varies, but the encoding always has a fixed size.

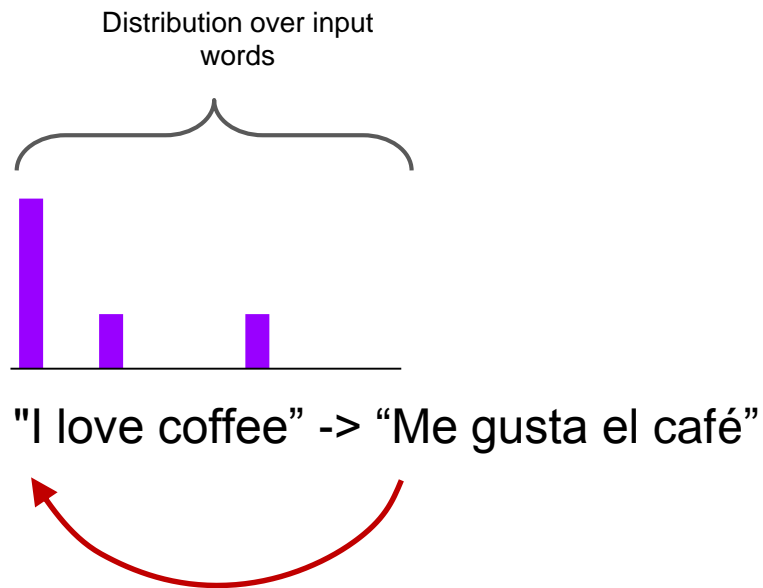
Soft Attention for Translation

“I love coffee” -> “Me gusta el café”


Bahdanau et al, “Neural Machine Translation by
Jointly Learning to Align and Translate”, ICLR 2015



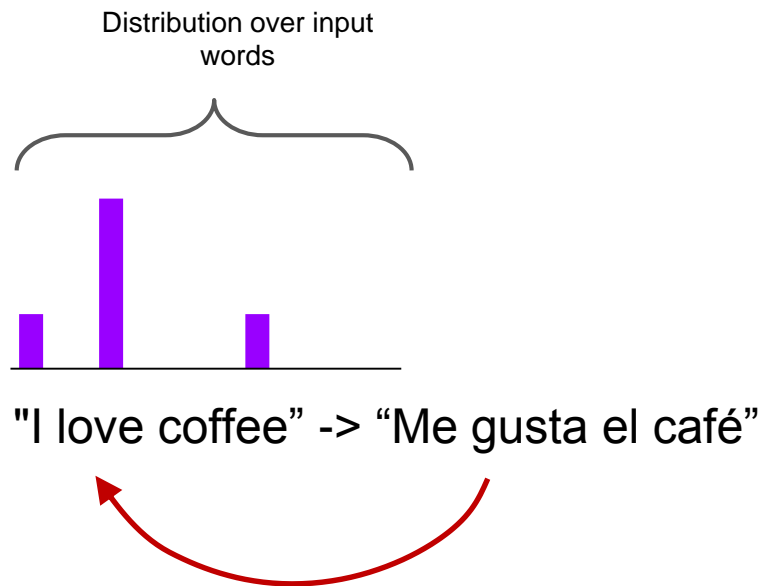
Soft Attention for Translation




Bahdanau et al, "Neural Machine Translation by Jointly Learning to Align and Translate", ICLR 2015




Soft Attention for Translation




Bahdanau et al, "Neural Machine Translation by Jointly Learning to Align and Translate", ICLR 2015




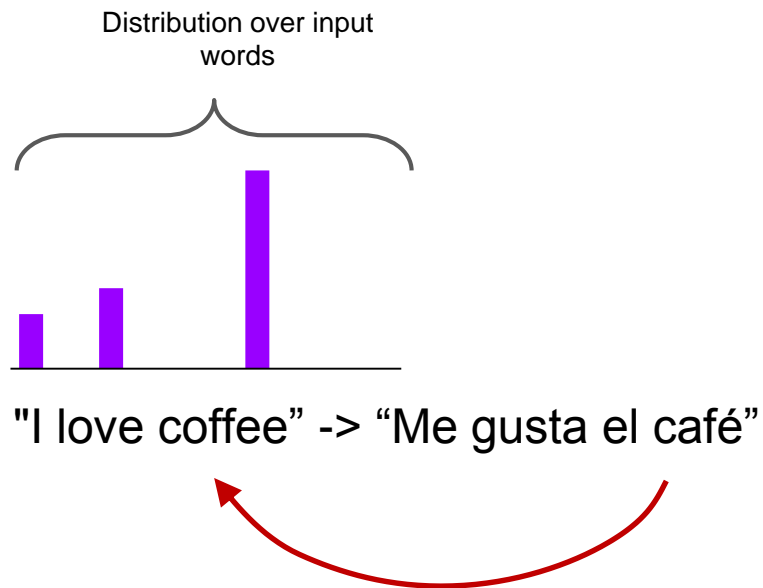
Soft Attention for Translation



Bahdanau et al, "Neural Machine Translation by Jointly Learning to Align and Translate", ICLR 2015




Soft Attention for Translation



Bahdanau et al, "Neural Machine Translation by
Jointly Learning to Align and Translate", ICLR 2015

Sequence-To-Sequence Criticisms


All the information from the source sentence has to pass through the bottleneck at the last unit(s) of the encoder.



Sentence length varies, but the encoding always has a fixed size.

Soft Attention for Translation – Bahdanau et al. model

For each output word, focus attention on a subset of all input words.



Context vector (input to decoder):
$$c_i = \sum_{j=1}^{T_x} \alpha_{ij} h_j$$

Mixture weights:


$$\alpha_{ij} = \frac{\exp(e_{ij})}{\sum_{k=1}^{T_x} \exp(e_{ik})}$$

Alignment score (how well do input words near j match output words at position i):

$$e_{ij} = a(s_{i-1}, h_j)$$


Encoder
(bidirectional RNN)

Soft Attention for Translation




From Y. Bengio CVPR 2015 Tutorial

Soft Attention for Translation




From Y. Bengio CVPR 2015 Tutorial

Soft Attention for Translation




From Y. Bengio CVPR 2015 Tutorial

Soft Attention for Translation




From Y. Bengio CVPR 2015 Tutorial

Soft Attention for Translation



(a)



(b)

Bahdanau et al, "Neural Machine Translation by Jointly Learning to Align and Translate", ICLR 2015

Soft Attention for Translation

Reached State of the art in one year:

(a) English→French (WMT-14)

	NMT(A)	Google	P-SMT
NMT	32.68	30.6*	37.03*
+Cand	33.28	—	
+UNK	33.99	32.7°	
+Ens	36.71	36.9°	

(b) English→German (WMT-15)

Model	Note
24.8	Neural MT
24.0	U.Edinburgh, Syntactic SMT
23.6	LIMS/KIT
22.8	U.Edinburgh, Phrase SMT
22.7	KIT, Phrase SMT

(c) English→Czech (WMT-15)


Model	Note
18.3	Neural MT
18.2	JHU, SMT+LM+OSM+Sparse
17.6	CU, Phrase SMT
17.4	U.Edinburgh, Phrase SMT
16.1	U.Edinburgh, Syntactic SMT

Criticism of Badanau et al.

The attention function $a(s_{i-1}, h_j)$ is rather complex, yet the attention often seems to be a simple heat map on word similarity:

The data path in Badanau is quite complicated: the attention path is another recurrent path between output states.


Doesn't generalize to deeper networks (shown to be Important by Sutskeyver et al.).



Luong and Manning added several architectural improvements.

Luong, Pham and Manning 2015


Stacked LSTM with arbitrary depth (c.f. bidirectional flat encoder in Bahdanau et al):



Global Attention Model



Global attention model is similar but simpler than Bahdanau's. It sits above the encoder/decoder and is not itself recurrent.

Different word matching functions were explored, some yielding better results.



Local Attention Model

- Compute a best aligned position p_t first
- Then compute a context vector centered at that position



Luong, Pham and Manning's Translation System (2015):

System	BLEU
Top – <i>NMT + 5-gram rerank</i> (Montreal)	24.9
Our ensemble 8 models + unk replace	25.9

Table 2: **WMT'15 English-German results** – *NIST* BLEU scores of the winning entry in WMT'15 and our best one on newstest2015.

System	Ppl.	BLEU
<i>WMT'15 systems</i>		
SOTA – <i>phrase-based</i> (Edinburgh)		29.2
NMT + 5-gram rerank (MILA)		27.6
<i>Our NMT systems</i>		
Base (reverse)	14.3	16.9
+ global (<i>location</i>)	12.7	19.1 (+2.2)
+ global (<i>location</i>) + feed	10.9	20.1 (+1.0)
+ global (<i>dot</i>) + drop + feed	9.7	22.8 (+2.7)
+ global (<i>dot</i>) + drop + feed + unk		24.9 (+2.1)

Table 3: **WMT'15 German-English results** –

Parsing

Recall (Lecture 10) RNNs ability to generate Latex, C code:

Proof. Omitted. □

Lemma 0.1. Let \mathcal{C} be a set of the construction.
 Let \mathcal{C} be a gerber covering. Let \mathcal{F} be a quasi-coherent sheaves of \mathcal{O} -modules. We have to show that

$$\mathcal{O}_{\mathcal{C}_X} = \mathcal{O}_X(\mathcal{C})$$

Proof. This is an algebraic space with the composition of sheaves \mathcal{F} on $X_{\text{étale}}$ we have

$$\mathcal{O}_X(\mathcal{F}) = \{ \text{morphisms } \times_{\mathcal{O}_X} (\mathcal{G}, \mathcal{F}) \}$$

where \mathcal{G} defines an isomorphism $\mathcal{F} \rightarrow \mathcal{F}$ of \mathcal{O} -modules. □

Lemma 0.2. This is an integer \mathbb{Z} is injective. □

Proof. See Spaces, Lemma ??.

Lemma 0.3. Let S be a scheme. Let X be a scheme and X is an affine open covering. Let $U \subset X$ be a canonical and locally of finite type. Let X be a scheme. Let X be a scheme which is equal to the formal complex.

The following to the construction of the lemma follows.

Let X be a scheme. Let X be a scheme covering. Let

$$b: X \rightarrow Y' \rightarrow Y \rightarrow Y' \times_X Y \rightarrow X.$$

be a morphism of algebraic spaces over S and Y .

Proof. Let X be a nonzero scheme of X . Let X be an algebraic space. Let \mathcal{F} be a quasi-coherent sheaf of \mathcal{O}_X -modules. The following are equivalent

- (1) \mathcal{F} is an algebraic space over S .
- (2) If X is an affine open covering.

Consider a common structure on X and X the functor $\mathcal{O}_X(U)$ which is locally of finite type. □

This since $\mathcal{F} \in \mathcal{F}$ and $x \in \mathcal{G}$ the diagram

is a limit. Then \mathcal{G} is a finite type and assume S is a flat and \mathcal{F} and \mathcal{G} is a \mathbb{Z} type \mathbb{Z} . This is of finite type diagrams, and

- the composition of \mathcal{G} is a regular sequence,
- \mathcal{O}_X is a sheaf of rings.

Proof. We have seen that $X = \text{Spec}(\mathbb{Z})$ and \mathcal{F} is a finite type representable algebraic space. The property \mathcal{F} is a finite morphism of algebraic stacks. The cohomology of X is an open neighbourhood of U .

Proof. This is clear that \mathcal{G} is a finite presentation, see Lemma ??.

A reduced scheme we conclude that U is an open covering of \mathcal{C} . The functor \mathcal{F} "field"

$$\mathcal{O}_{\mathcal{C}_X} \rightarrow \mathcal{F}_X \rightarrow \mathcal{O}_{\mathcal{C}_{\text{étale}}} \rightarrow \mathcal{O}_{\mathcal{C}_X}^{\text{étale}}(\mathcal{O}_{\mathcal{C}_X}^{\text{étale}})$$

is an isomorphism of covering of \mathcal{O}_X . If \mathcal{F} is the unique element of \mathcal{F} such this is an isomorphism.

The property \mathcal{F} is a disjoint union of Proposition ?? and we can filtered a presentations of a scheme \mathcal{O}_X -algebra with \mathcal{F} are opens of finite type over S . If \mathcal{F} is a scheme theoretic image points.

If \mathcal{F} is a finite direct sum \mathcal{O}_X is a closed immersion, see Lemma ??.

This sequence of \mathcal{F} is a similar morphism.


```
static void do_command(struct seq_file *m, void *v)
{
    int column = 32 << (cmd[2] & 0x80);
    if (state)
        cmd = (int)(int_state ^ (in_8(&ch->ch_flags) & Cmd) ? 2 : 1);
    else
        seq = 1;
    for (i = 0; i < 16; i++) {
        if (k & (1 << 1))
            pipe = (in_use & MMTHREAD_UNCCA) +
                ((count & 0x00000000ffffffff) & 0x0000000f) << 8;
        if (count == 0)
            sub(pid, ppc_md.kexec_handle, 0x20000000);
        pipe_set_bytes(i, 0);
    }
    /* Free our user pages pointer to place camera if all dash */
    subsystem_info = &of_changes[PAGE_SIZE];
    rek_controls(offset, idx, &offset);
    /* Now we want to deliberately put it to device */
    control_check_polarity(&context, val, 0);
    for (i = 0; i < COUNTER; i++)
        seq_puts(s, "policy ");
}
```

They seem to do well with tree-structured data.

What about natural language parsing?


Parsing

Sequence models generate linear structures, but these can easily encode trees by “closing parens” (prefix tree notation):



John has a dog . → (S (NP NNP)_{NP} (VP VBZ (NP DT NN)_{NP})_{VP} .)_S

Parsing Cheat Sheet



John has a dog . → (S (NP NNP)_{NP} (VP VBZ (NP DT NN)_{NP})_{VP} .)_S

S = Sentence

NP = Noun Phrase

VP = Verb Phrase

NNP = Proper Noun ("John")


VBZ = Verb, 3rd person, singular ("has")

DT = Determiner ("a")

NN = Noun, singular ("dog")

A Sequence-To-Sequence Parser

The model is a depth-3 sequence-to-sequence predictor, augmented with the attention model of Bahdanau 2014.



Grammar as a Foreign Language Oriol Vinyals, Google, Lukasz Kaiser, Terry Koo, Slav Petrov, Ilya Sutskever, Geoffrey Hinton, NIPS 2015

“Neural machine translation by jointly learning to align and translate.” Dzmitry Bahdanau, Kyunghyun Cho, and Yoshua Bengio. arXiv 2014.

A Sequence-To-Sequence Parser

Chronology:

- First tried training a basic sequence-to-sequence model on human-annotated training treebanks. **Poor results.**
- Then training on parse trees **generated by the Berkeley Parser**, achieved similar performance (90.5 F1 score) to it.
- Next added the attention model, trained **on human treebank data**, also achieved 90.5 F1.
- Finally, created a synthetic dataset of **high-confidence parse trees** (agreed on by two parsers). Achieved a new state-of-the-art of 92.5 F1 score (WSJ dataset).

F1 is a widely-used accuracy measure that combines precision and recall

A Sequence-To-Sequence Parser

Quick Training Details:

- Depth = 3, layer dimension = 256.
- **Dropout** between layers 1 and 2, and 2 and 3.
- **No POS tags!!** Improved by F1 1 point by leaving them out.
- Input reversing.

Attention-only Translation Models

Problems with recurrent networks:

- **Sequential training and inference**: time grows in proportion to sentence length. Hard to parallelize.
- **Long-range dependencies** have to be remembered across many single time steps.
- **Tricky to learn hierarchical structures** (“car”, “blue car”, “into the blue car”...)


Alternative:

- Convolution – but has other limitations.

Self-Attention

Information flows from within the same subnetwork (either encoder or decoder).
Convolution applies fixed transform weights. Self-attention applies variable weights
(but typically not transformations):

Convolution



Self-Attention





image from Lukas Kaiser, Stanford NLP seminar

Self-Attention “Transformers” (not spatial transformers)

- Constant path length between any two positions.
- Variable receptive field (or the whole input sequence).
- Supports hierarchical information flow by stacking self-attention layers.
- Trivial to parallelize.
- Attention weighting controls information propagation.
- **Can replace word-based recurrence entirely.**

Vaswani et al. “Attention is all you need”, arXiv 2017

Attention in Transformer Networks




We saw this in Bahdanau and Luong models

Encoder-Decoder Attention




Encoder Self-Attention



MaskedDecoder Self-Attention

image from Lukas Kaiser, Stanford NLP seminar

Attention in Transformer Networks




Encoder-Decoder Attention

Replaces word recurrence in
encoder and decoder




Encoder Self-Attention



MaskedDecoder Self-Attention

image from Lukas Kaiser, Stanford NLP seminar


Attention in Transformer Networks



Encoder-Decoder Attention



Encoder Self-Attention



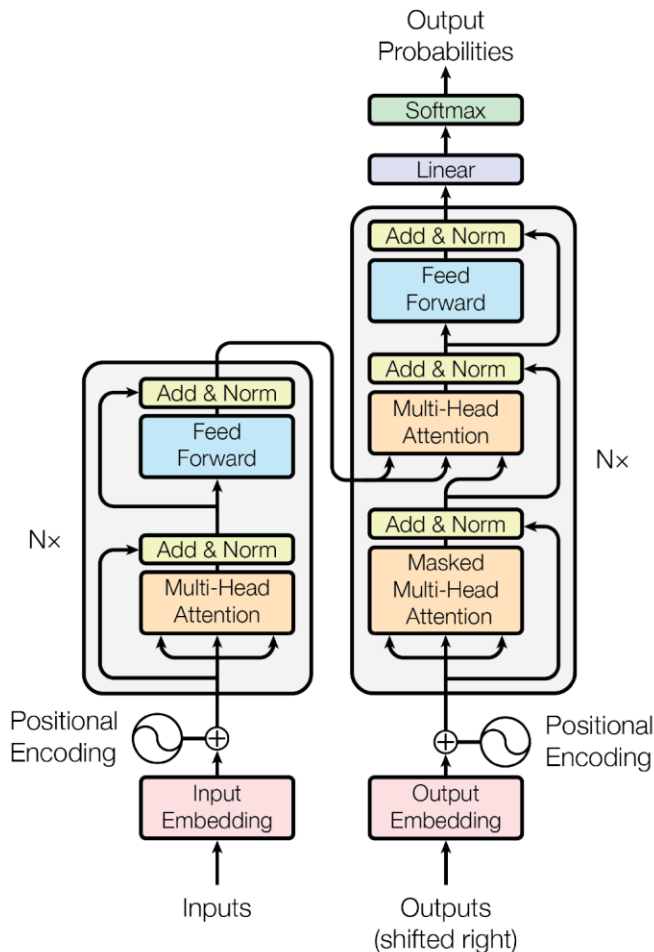
Masked Decoder Self-Attention

Masking limits attention to earlier units:
 y_i depends only on y_j for $j < i$.

image from Lukas Kaiser, Stanford NLP seminar

The Transformer

- Basic unit shown at right.
- In experiments, stacked with $N=6$.
- Output words fed back as input, shifted right. Can use beam search as before.
- Inputs and outputs are embedded in vector spaces of fixed dimension.
- Positional encoding: when words are combined through attention, their location is lost. Positional encoding adds it back.



Attention Implementation

- Attention is modeled as a key-value store:

Q = query vector

K = key


V = value

Encoder-decoder layer: the queries come from the previous decoder layer, and the memory keys and values come from the output of the encoder. (Similar to Bahdanau).

Self-attention layer: all of the keys, values and queries come from the output of the previous layer in the encoder.

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$


Scaled Dot-Product Attention




Multi-Headed Attention

- Simple attention blends the results of all the attended-to inputs. It doesn't allow a per-input transformation, as convolution does.
- The solution is to use “multi-headed attention”:


Convolution




Multi-Head Attention



Multi-Headed Attention




Multi-Headed Attention



Anaphora (pronoun or article) resolution

Multi-Headed Attention



Anaphora (pronoun or article) resolution

Transformer Results

Machine Translation Results: WMT-14

Model	BLEU		Training Cost (FLOPs)	
	EN-DE	EN-FR	EN-DE	EN-FR
ByteNet [17]	23.75			
Deep-Att + PosUnk [37]		39.2		$1.0 \cdot 10^{20}$
GNMT + RL [36]	24.6	39.92	$2.3 \cdot 10^{19}$	$1.4 \cdot 10^{20}$
ConvS2S [9]	25.16	40.46	$9.6 \cdot 10^{18}$	$1.5 \cdot 10^{20}$
MoE [31]	26.03	40.56	$2.0 \cdot 10^{19}$	$1.2 \cdot 10^{20}$
Deep-Att + PosUnk Ensemble [37]		40.4		$8.0 \cdot 10^{20}$
GNMT + RL Ensemble [36]	26.30	41.16	$1.8 \cdot 10^{20}$	$1.1 \cdot 10^{21}$
ConvS2S Ensemble [9]	26.36	41.29	$7.7 \cdot 10^{19}$	$1.2 \cdot 10^{21}$
Transformer (base model)	27.3	38.1	$3.3 \cdot 10^{18}$	
Transformer (big)	28.4	41.0	$2.3 \cdot 10^{19}$	

English-to-English Translation ?!

Yes, it does make sense. a.k.a. summarization.


Liu et al, "GENERATING WIKIPEDIA BY SUMMARIZING LONG SEQUENCES" arXiv 2018

M = input length, N = output length


Summarization: $M \gg N$



Encoder-Decoder Attention



Encoder Self-Attention



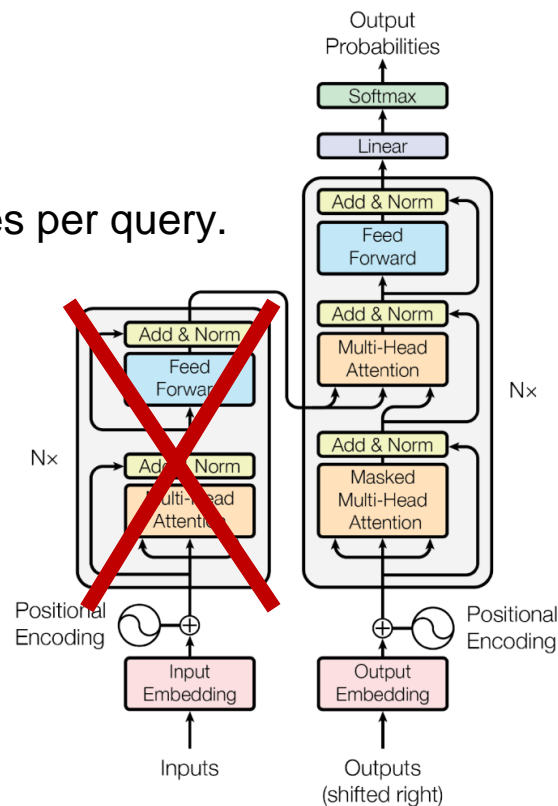
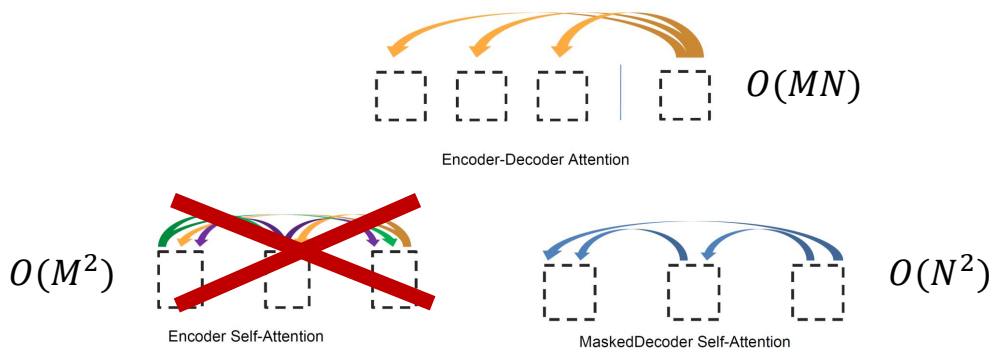
MaskedDecoder Self-Attention

Large-scale Summarization (Wikipedia)

Like translation, but we completely remove the encoder.

Source data (large!):

- The references for a Wikipedia article.
- Web search using article section titles, ~ 10 web pages per query.



Large-scale Summarization

Results:

Model	Test perplexity	ROUGE-L
<i>seq2seq-attention, $L = 500$</i>	5.04952	12.7
<i>Transformer-ED, $L = 500$</i>	2.46645	34.2
<i>Transformer-D, $L = 4000$</i>	2.22216	33.6
<i>Transformer-DMCA, no MoE-layer, $L = 11000$</i>	2.05159	36.2
<i>Transformer-DMCA, MoE-128, $L = 11000$</i>	1.92871	37.9
<i>Transformer-DMCA, MoE-256, $L = 7500$</i>	1.90325	38.8

L = input window length.

ED = encoder-decoder.

D = decoder only.

DMCA = a memory compression technique (strided convolution).

MoE = mixture of experts layer.

Translation Takeaways

- Sequence-to-sequence translation
 - Input reversal
 - Narrow beam search
- Adding Attention
 - Compare latent states of encoder/decoder (Bahdanau).
 - Simplify and avoid more recurrence (Luong).



Translation Takeaways

- Parsing as translation:
 - Translation models can solve many “transduction” tasks.
- Attention only models:
 - Self-attention replaces recurrence, improves performance.
 - Use depth to model hierarchical structure.
 - Multi-headed attention allows interpretation of inputs.

