

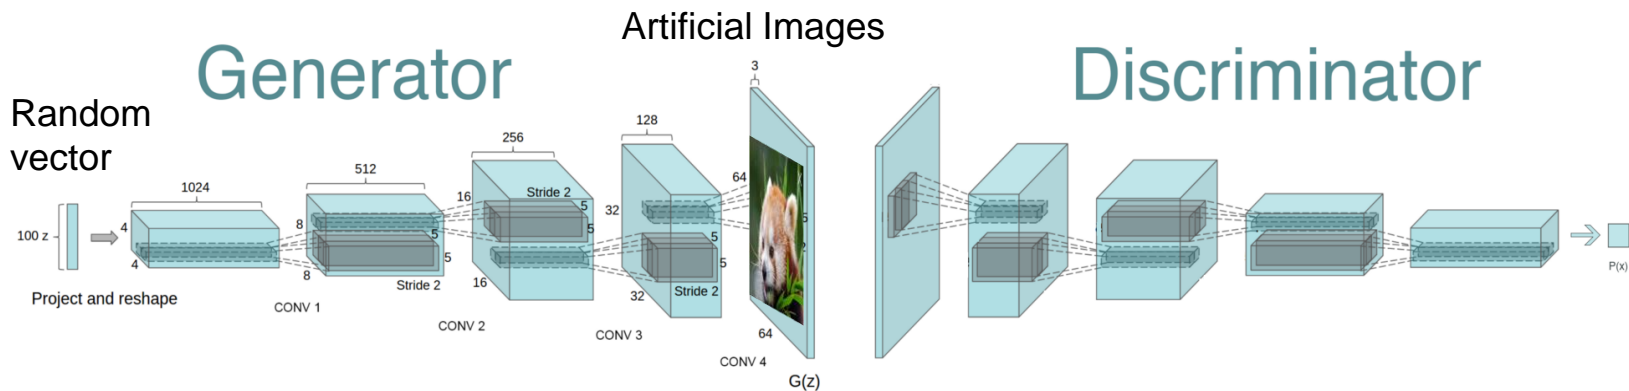
CS194/294-129: Designing, Visualizing and Understanding Deep Neural Networks

John Canny

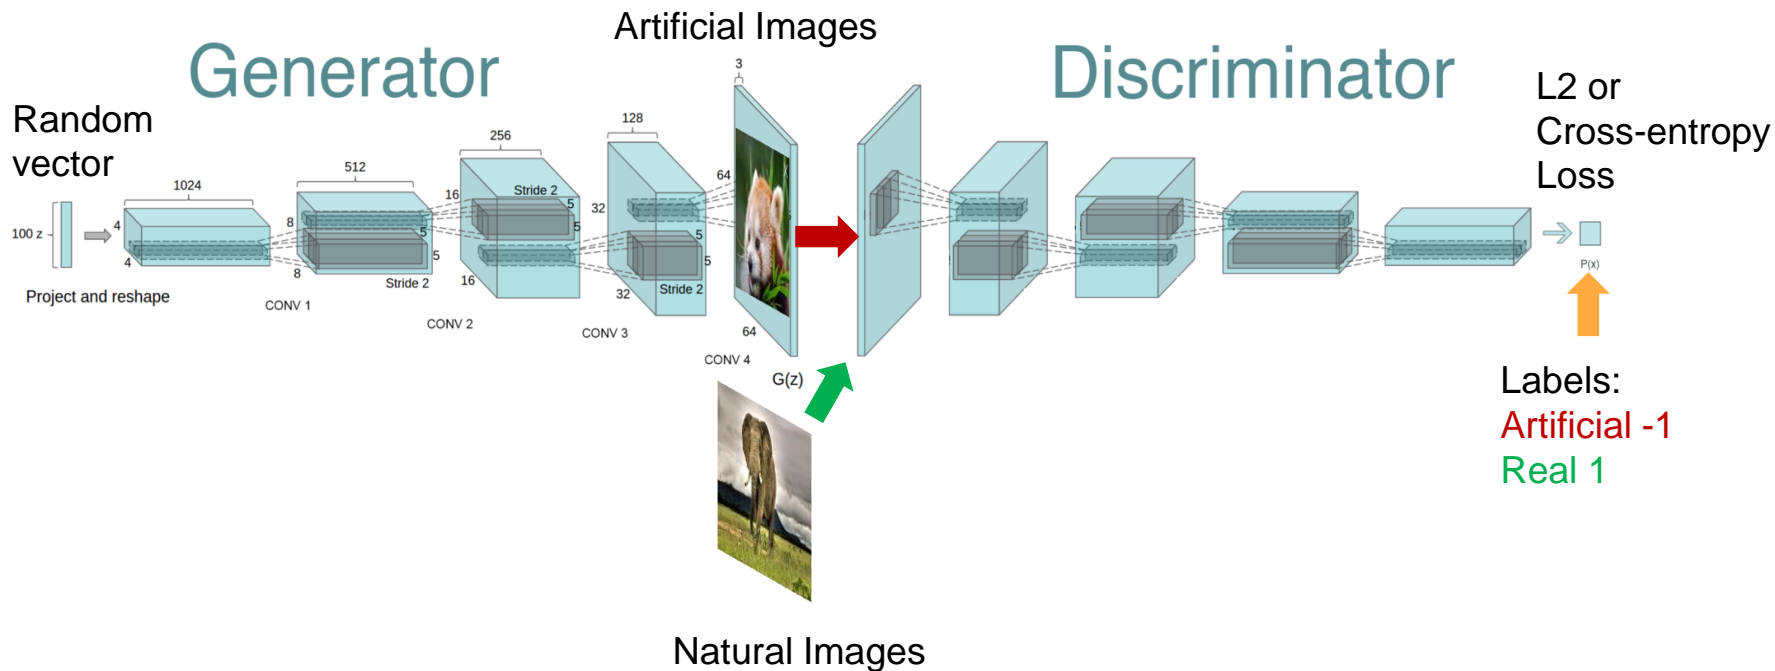
Spring 2018

Lecture 17: Imitation Learning

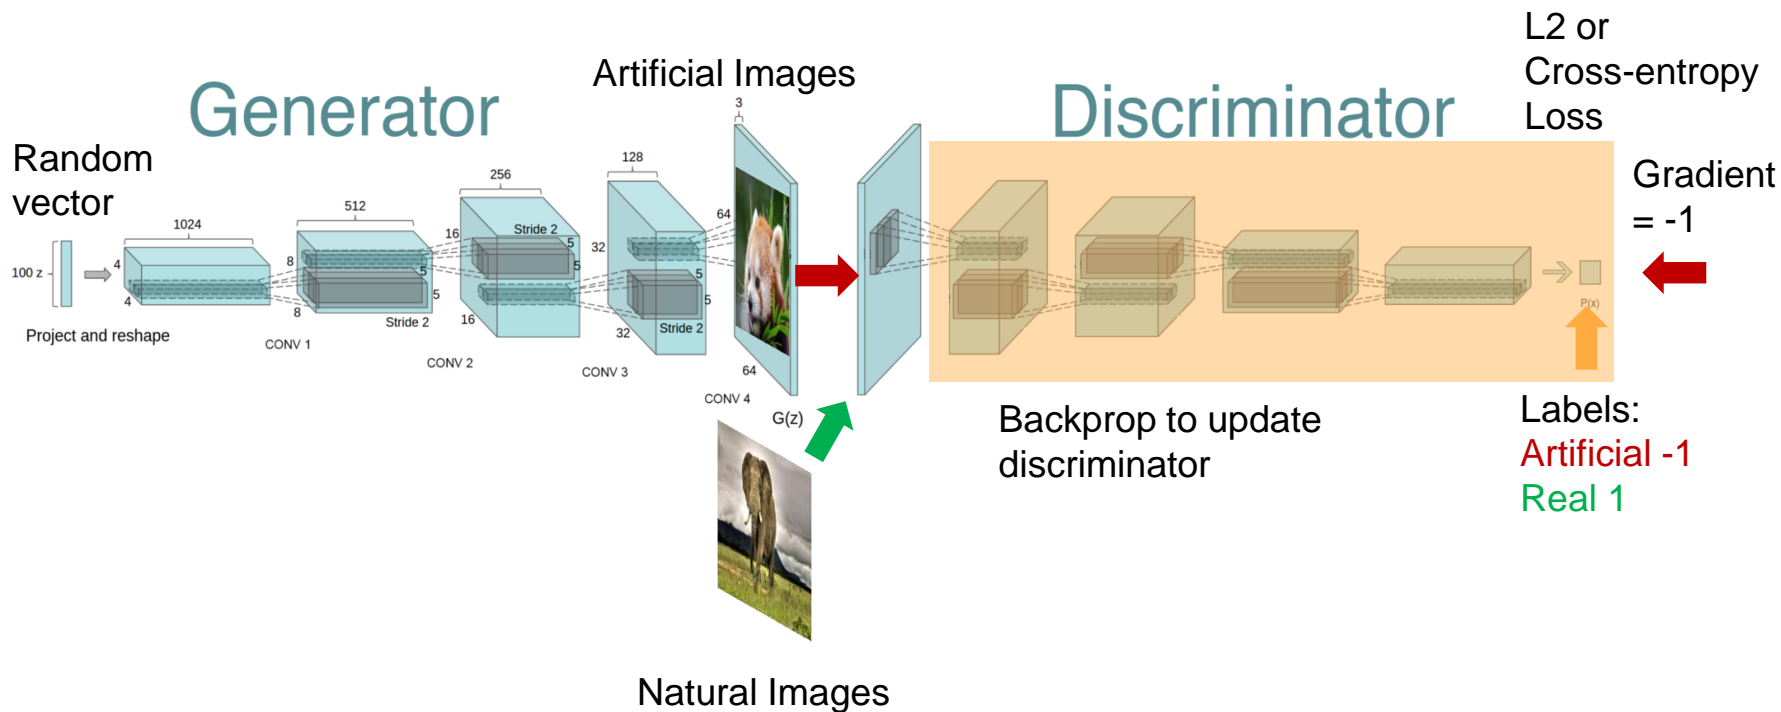
Last Time: Generative Adversarial Networks (GANs)



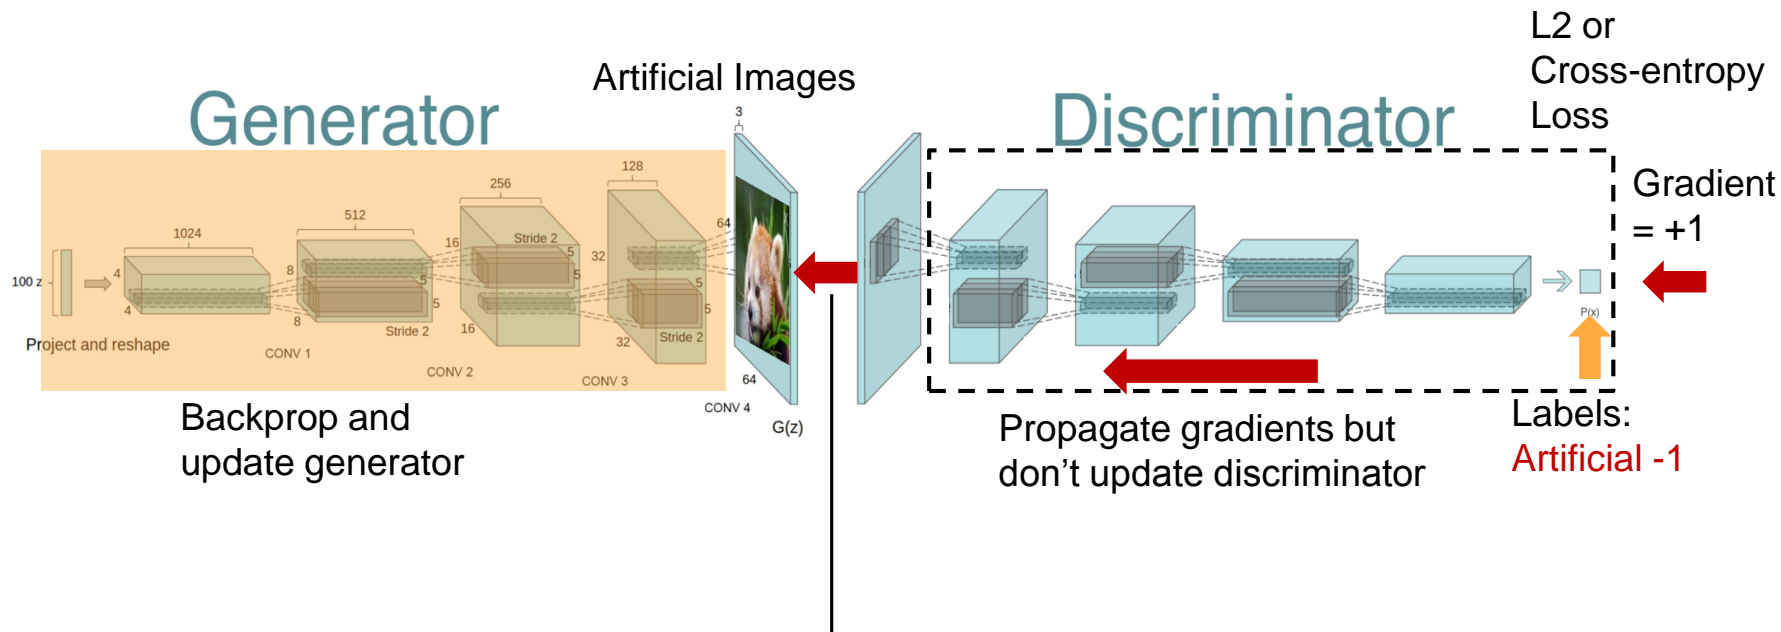
Last Time: GAN Discriminator Training



GAN Training: Minimize Discriminator classification loss

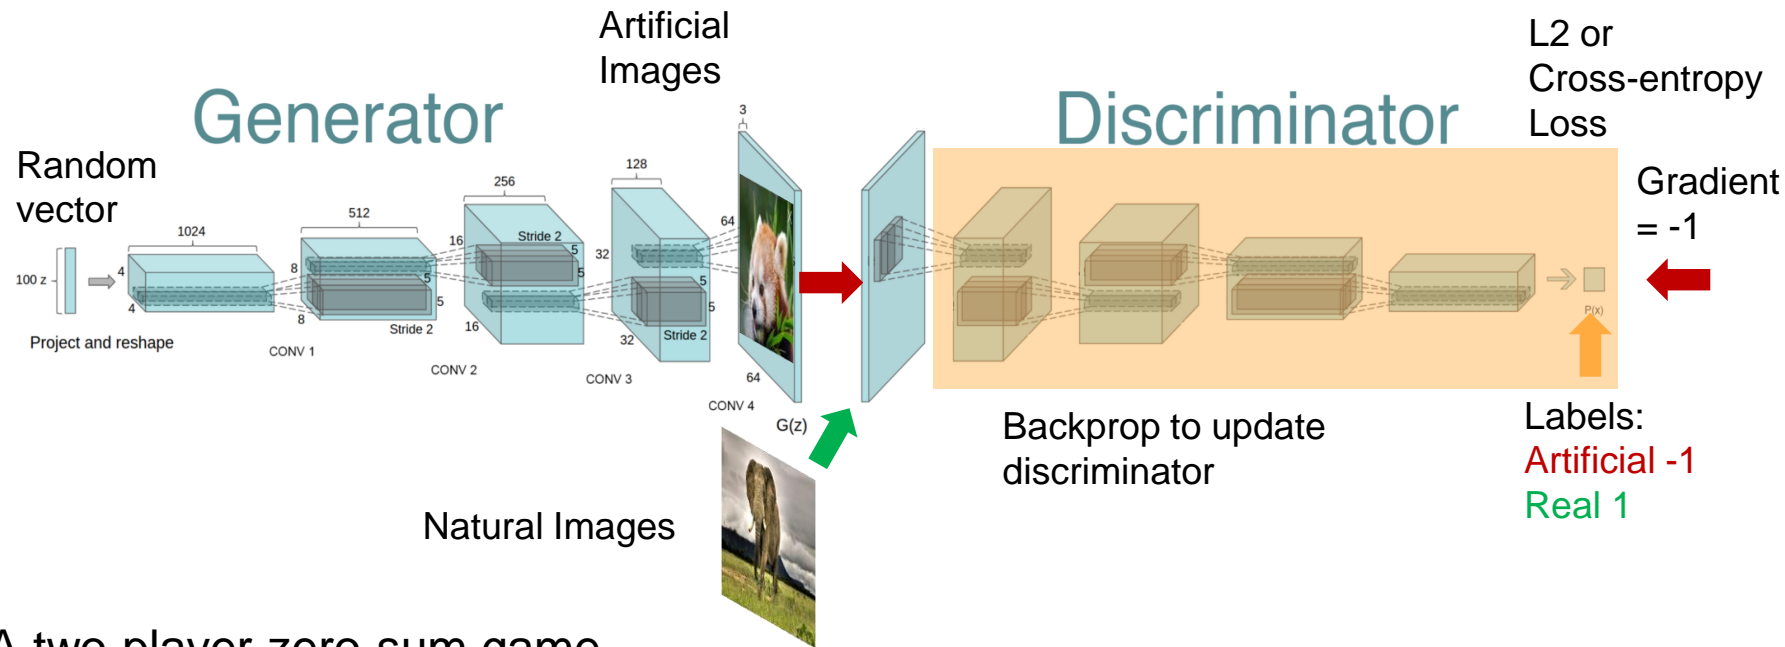


GAN Training: Train Generator to Fool the Discriminator

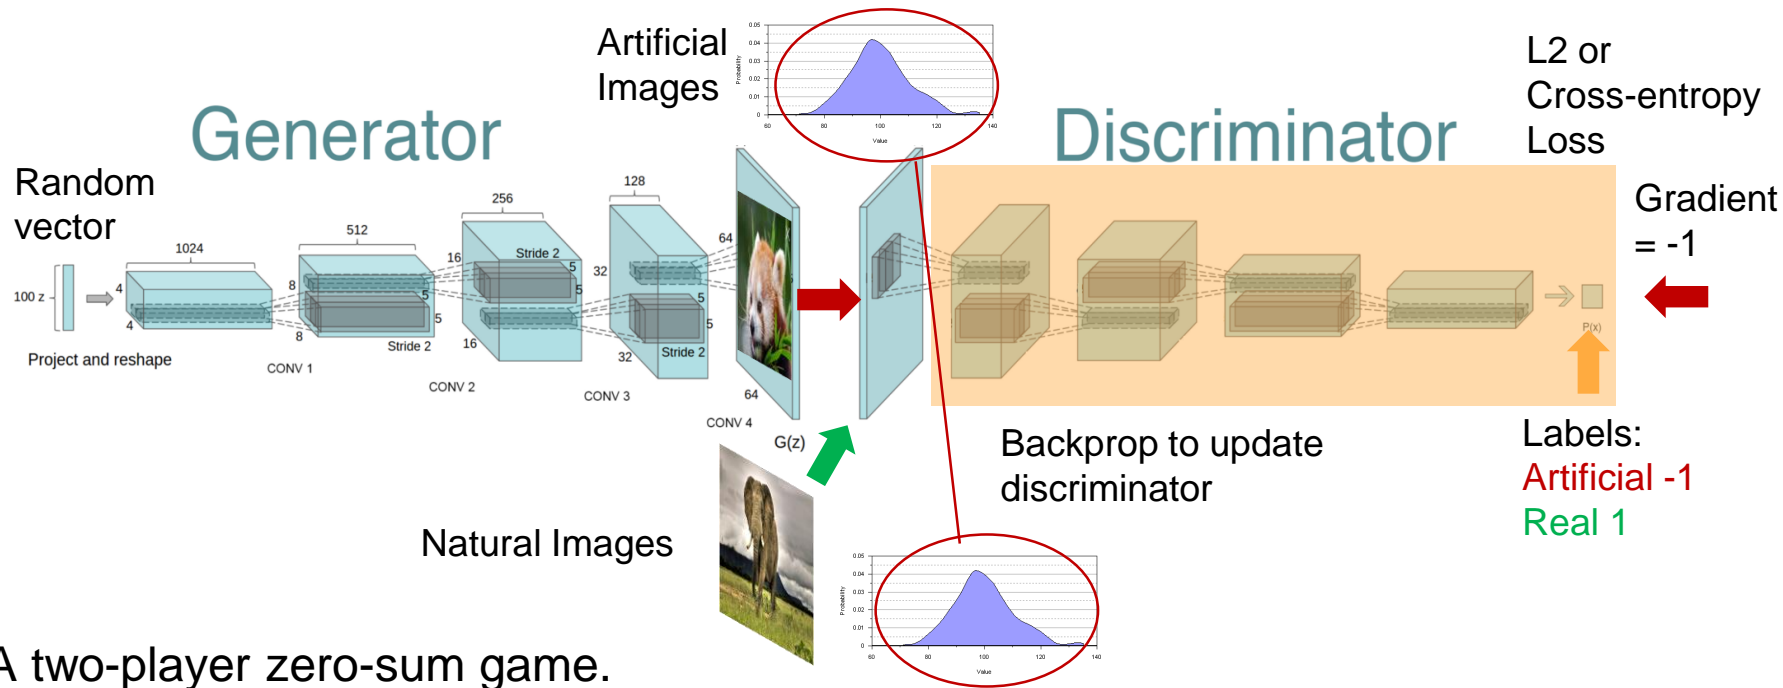


This gradient nudges the image from "artificial" toward "natural"

GAN Training: Alternate Discriminator/Generator training




GAN Training: Alternate Discriminator/Generator training



A two-player zero-sum game.

Optimizing with minimax (alternating optimization) minimizes the difference (Jensen-Shannon divergence) between generator and natural image probability distributions.

This Time: Deep Control



Deep Control: First Idea: Imitate Human Actions

Supervised training of deep networks (with image category labels, captions, translations,...) from human data has worked well so far...

What about mimicking human control actions?




Imitation Learning via Behavior Cloning

This approach is called **behavior cloning**. Note that its not enough to record human actions, because humans are constantly adapting to the world.

We need to learn a **control loop** from sensors to actuators.




Sensorimotor Learning




tiger
tiger
tiger cat
jaguar
lynx


Sensorimotor Learning



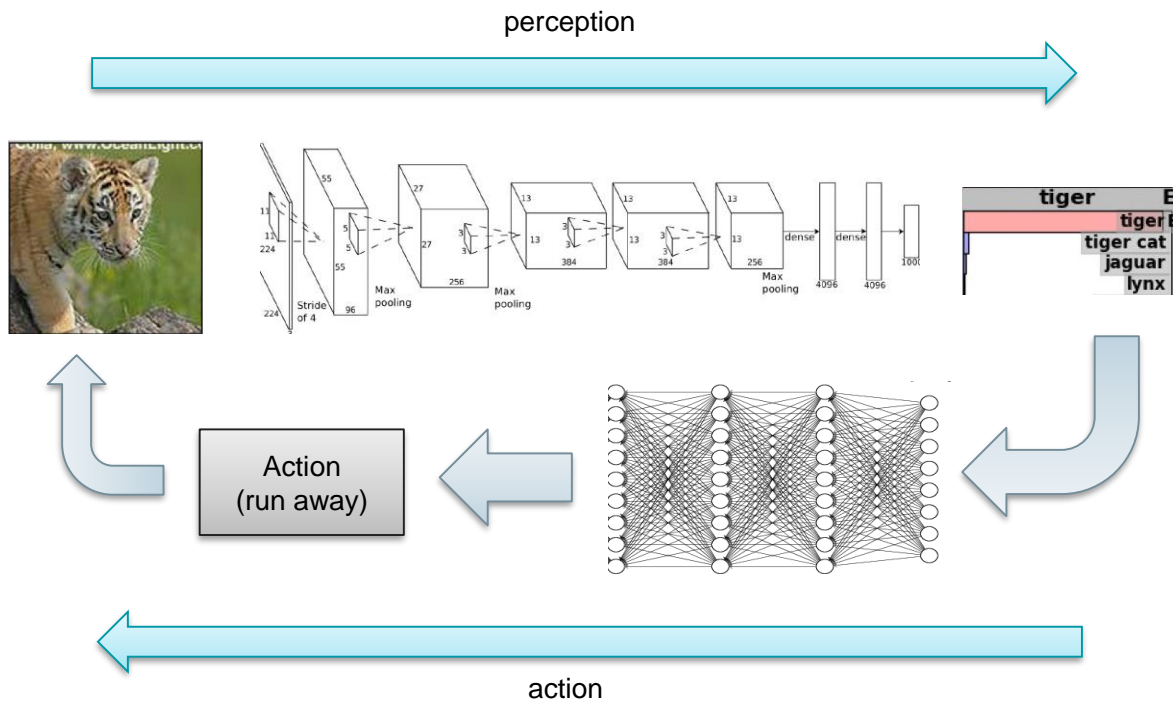
Sensorimotor Learning




Sensorimotor Learning




Sensorimotor Learning



sensorimotor loop




Action
(run away)




End-to-end vision

standard
computer
vision

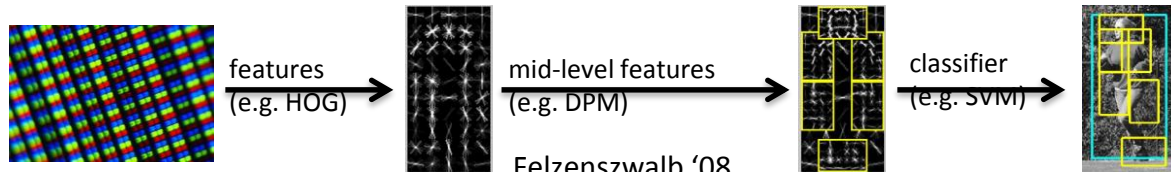


deep
learning




End-to-end vision

standard
computer
vision




deep
learning



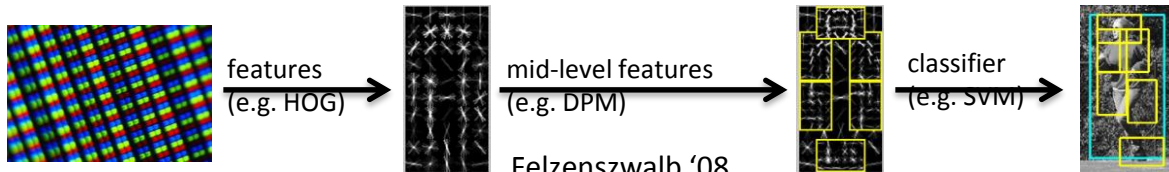
End-to-end control

standard
robotic
control




End-to-end vision

standard
computer
vision




deep
learning






End-to-end control



standard
robotic
control

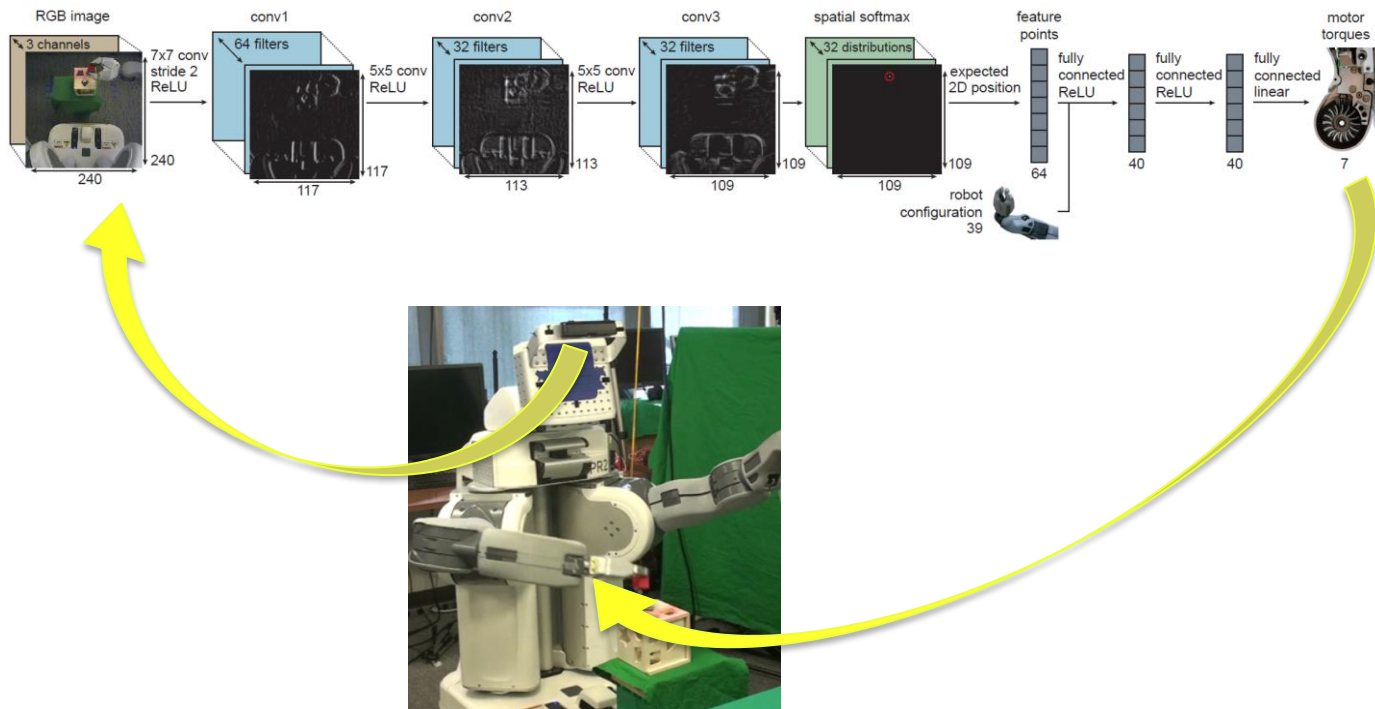


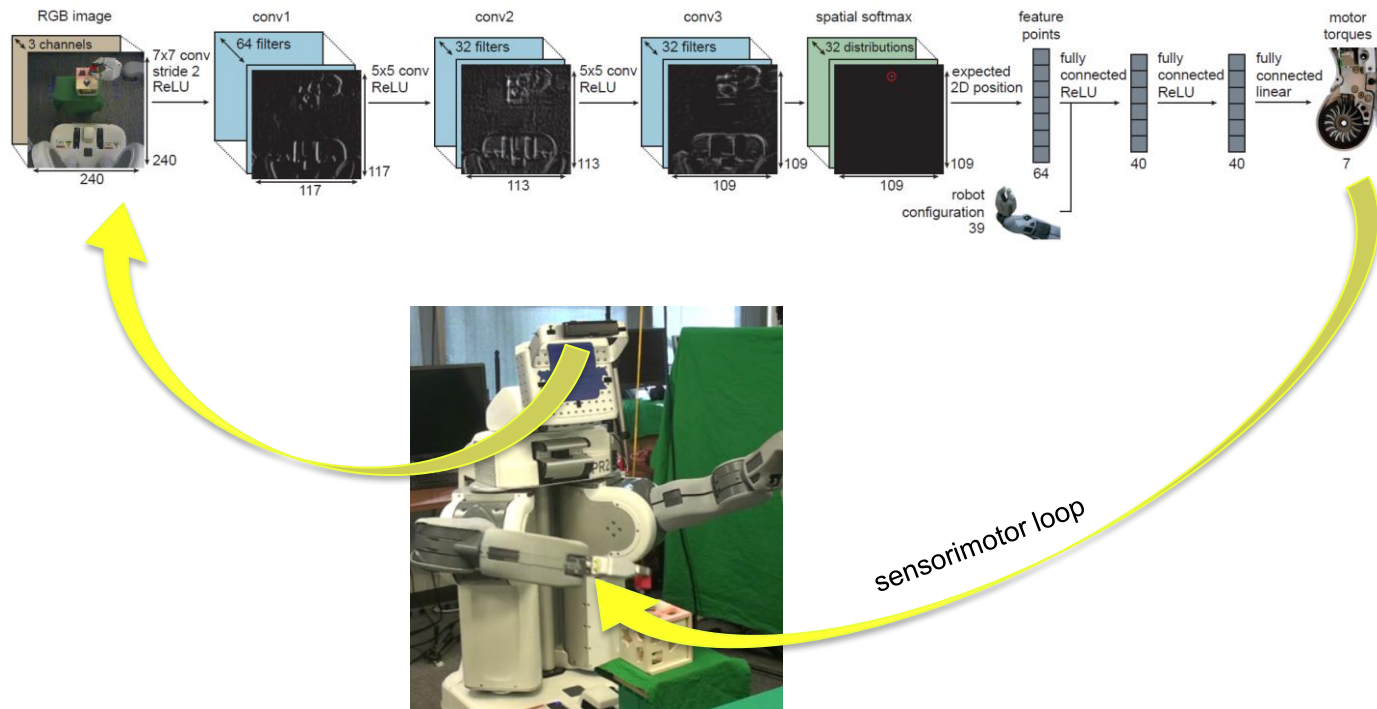
deep
sensorimotor
learning

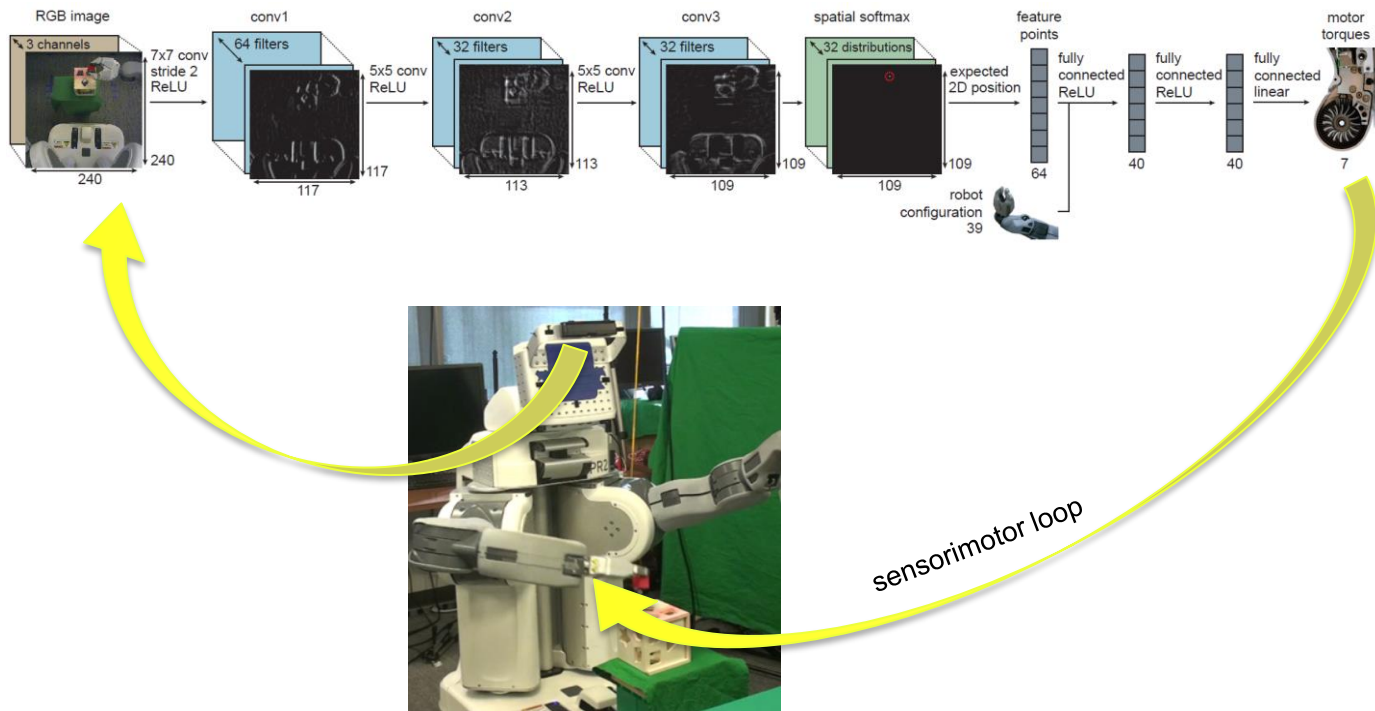




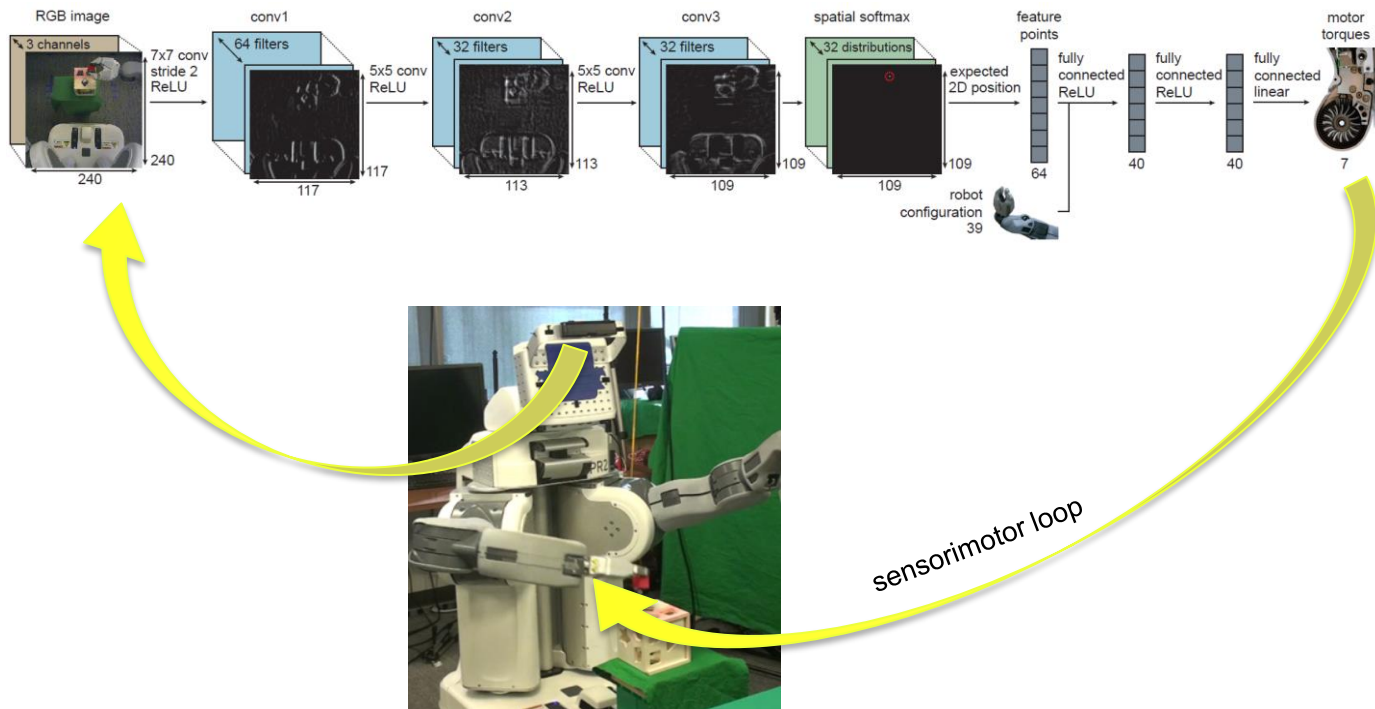








indirect supervision




indirect supervision
actions have consequences


Terminology & notation



\mathbf{o}



$$\pi_{\theta}(\mathbf{u}|\mathbf{o})$$




\mathbf{u}


Terminology & notation



\mathbf{o}_t




$\pi_{\theta}(\mathbf{u}_t | \mathbf{o}_t)$




\mathbf{u}_t


Terminology & notation




Terminology & notation




Terminology & notation




Terminology & notation




Terminology & notation



Terminology & notation




Terminology & notation



\mathbf{o}_t – observation

\mathbf{u}_t – action

Terminology & notation




\mathbf{o}_t – observation

\mathbf{u}_t – action

$\pi_\theta(\mathbf{u}_t | \mathbf{o}_t)$ – policy

Terminology & notation




\mathbf{x}_t – state

\mathbf{o}_t – observation

\mathbf{u}_t – action

$\pi_\theta(\mathbf{u}_t | \mathbf{o}_t)$ – policy

Terminology & notation



\mathbf{x}_t – state


\mathbf{o}_t – observation

\mathbf{u}_t – action

$\pi_{\theta}(\mathbf{u}_t | \mathbf{o}_t)$ – policy



Terminology & notation



\mathbf{x}_t – state

\mathbf{o}_t – observation


\mathbf{u}_t – action

$\pi_\theta(\mathbf{u}_t | \mathbf{o}_t)$ – policy



\mathbf{o}_t – observation

Terminology & notation




\mathbf{x}_t – state

\mathbf{o}_t – observation


\mathbf{u}_t – action

$\pi_\theta(\mathbf{u}_t | \mathbf{o}_t)$ – policy



\mathbf{o}_t – observation

Terminology & notation



\mathbf{x}_t – state

\mathbf{o}_t – observation

\mathbf{u}_t – action

$\pi_{\theta}(\mathbf{u}_t | \mathbf{o}_t)$ – policy




\mathbf{o}_t – observation



\mathbf{x}_t – state

Terminology & notation



\mathbf{x}_t – state

\mathbf{o}_t – observation

\mathbf{u}_t – action

$\pi_\theta(\mathbf{u}_t | \mathbf{o}_t)$ – policy




\mathbf{o}_t – observation




\mathbf{x}_t – state

Terminology & notation



Terminology & notation




\mathbf{x}_t – state


\mathbf{o}_t – observation

\mathbf{u}_t – action

$\pi_{\theta}(\mathbf{u}_t | \mathbf{o}_t)$ – policy



Terminology & notation




\mathbf{x}_t – state


\mathbf{o}_t – observation

\mathbf{u}_t – action

$\pi_\theta(\mathbf{u}_t | \mathbf{o}_t)$ – policy



Terminology & notation




\mathbf{x}_t – state


\mathbf{o}_t – observation

\mathbf{u}_t – action

$\pi_\theta(\mathbf{u}_t | \mathbf{o}_t)$ – policy



Terminology & notation




\mathbf{x}_t – state


\mathbf{o}_t – observation

\mathbf{u}_t – action

$\pi_{\theta}(\mathbf{u}_t | \mathbf{o}_t)$ – policy



Terminology & notation




\mathbf{x}_t – state


\mathbf{o}_t – observation

\mathbf{u}_t – action

$\pi_{\theta}(\mathbf{u}_t | \mathbf{o}_t)$ – policy



Terminology & notation



\mathbf{x}_t – state

\mathbf{o}_t – observation

\mathbf{u}_t – action


$\pi_{\theta}(\mathbf{u}_t | \mathbf{o}_t)$ – policy

a bit of history...

\mathbf{x}_t – state

\mathbf{u}_t – action

Terminology & notation



\mathbf{x}_t – state

\mathbf{o}_t – observation

\mathbf{u}_t – action

$\pi_{\theta}(\mathbf{u}_t | \mathbf{o}_t)$ – policy

a bit of history...


\mathbf{x}_t – state

\mathbf{u}_t – action

\mathbf{s}_t – state

\mathbf{a}_t – action

Terminology & notation



\mathbf{x}_t – state

\mathbf{o}_t – observation

\mathbf{u}_t – action

$\pi_{\theta}(\mathbf{u}_t | \mathbf{o}_t)$ – policy

a bit of history...

\mathbf{x}_t – state

\mathbf{u}_t – action




Lev Pontryagin

\mathbf{s}_t – state

\mathbf{a}_t – action

Terminology & notation



\mathbf{x}_t – state

\mathbf{o}_t – observation

\mathbf{u}_t – action

$\pi_{\theta}(\mathbf{u}_t | \mathbf{o}_t)$ – policy

a bit of history...

\mathbf{x}_t – state

\mathbf{u}_t – action



Lev Pontryagin


\mathbf{s}_t – state

\mathbf{a}_t – action



Richard Bellman

Terminology & notation



\mathbf{x}_t – state

\mathbf{o}_t – observation

\mathbf{u}_t – action

$\pi_{\theta}(\mathbf{u}_t | \mathbf{o}_t)$ – policy

a bit of history...

\mathbf{x}_t – state

\mathbf{u}_t – action

управление



Lev Pontryagin


\mathbf{s}_t – state

\mathbf{a}_t – action



Richard Bellman

Terminology & notation



\mathbf{x}_t – state

\mathbf{o}_t – observation

\mathbf{u}_t – action

$\pi_{\theta}(\mathbf{u}_t | \mathbf{o}_t)$ – policy

a bit of history...

\mathbf{x}_t – state

\mathbf{u}_t – action

управление



Lev Pontryagin



\mathbf{s}_t – state

\mathbf{a}_t – action




Richard Bellman


Imitation Learning



\mathbf{o}_t




$\pi_{\theta}(\mathbf{u}_t | \mathbf{o}_t)$





\mathbf{u}_t



Imitation Learning



Imitation Learning





Imitation Learning





\mathbf{o}_t

\mathbf{u}_t



Imitation Learning



Imitation Learning



Imitation Learning




Does it work?

No!

If the training trajectory error is ϵ , Ross et al. show that errors in the learned model's trajectory can be order $O(T^2\epsilon)$

Does it work?


No!



If the training trajectory error is ϵ , Ross et al. show that errors in the learned model's trajectory can be order $O(T^2\epsilon)$

Does it work?


No!



If the training trajectory error is ϵ , Ross et al. show that errors in the learned model's trajectory can be order $O(T^2\epsilon)$

Does it work?


No!



If the training trajectory error is ϵ , Ross et al. show that errors in the learned model's trajectory can be order $O(T^2\epsilon)$

Does it work?


No!



If the training trajectory error is ϵ , Ross et al. show that errors in the learned model's trajectory can be order $O(T^2\epsilon)$

Does it work?


No!



If the training trajectory error is ϵ , Ross et al. show that errors in the learned model's trajectory can be order $O(T^2\epsilon)$

Does it work?

No!



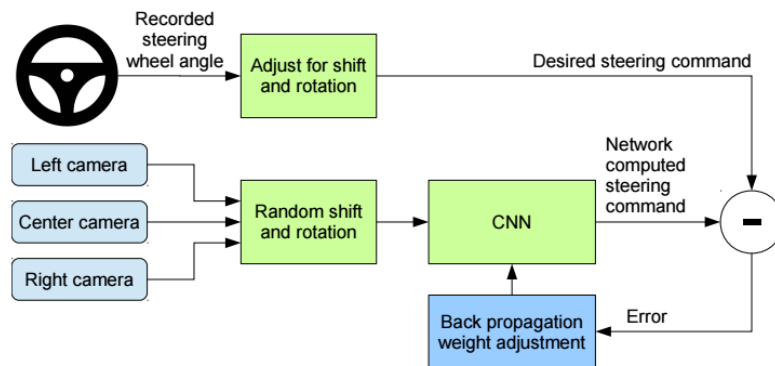
If the training trajectory error is ϵ , Ross et al. show that errors in the learned model's trajectory can be order $O(T^2\epsilon)$

Does it work?

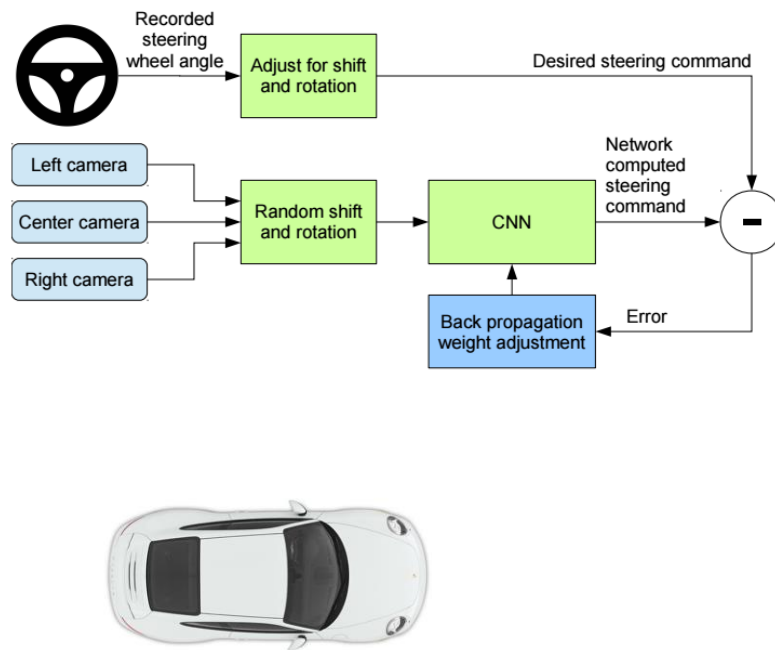
Yes!



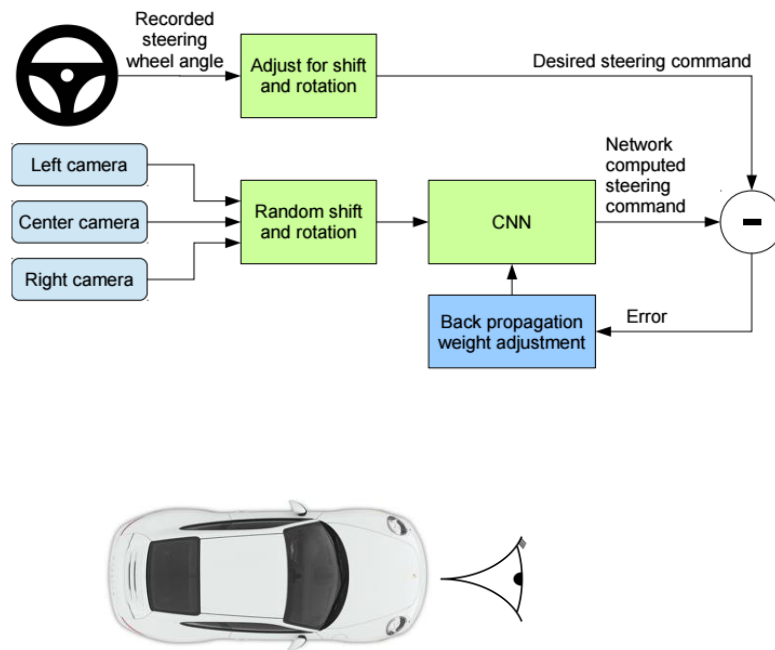
Why did that work?



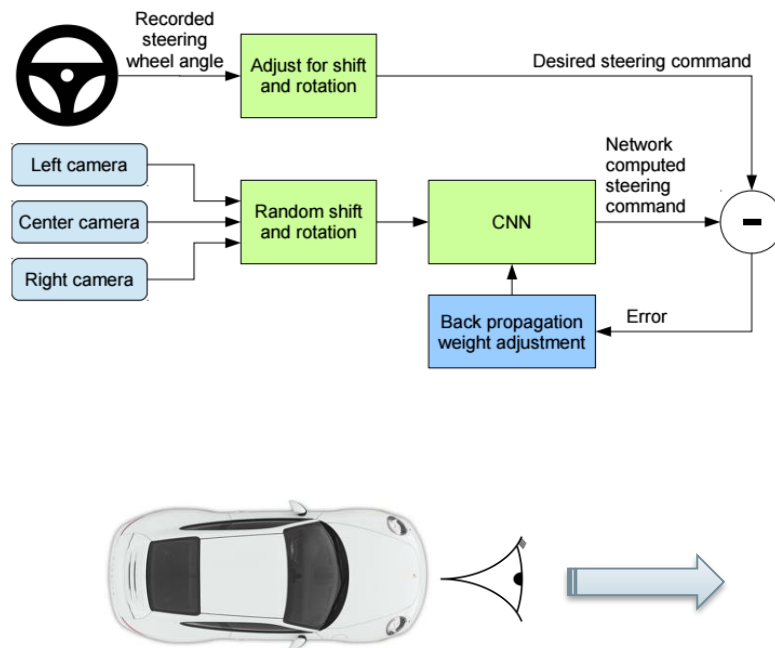
Why did that work?



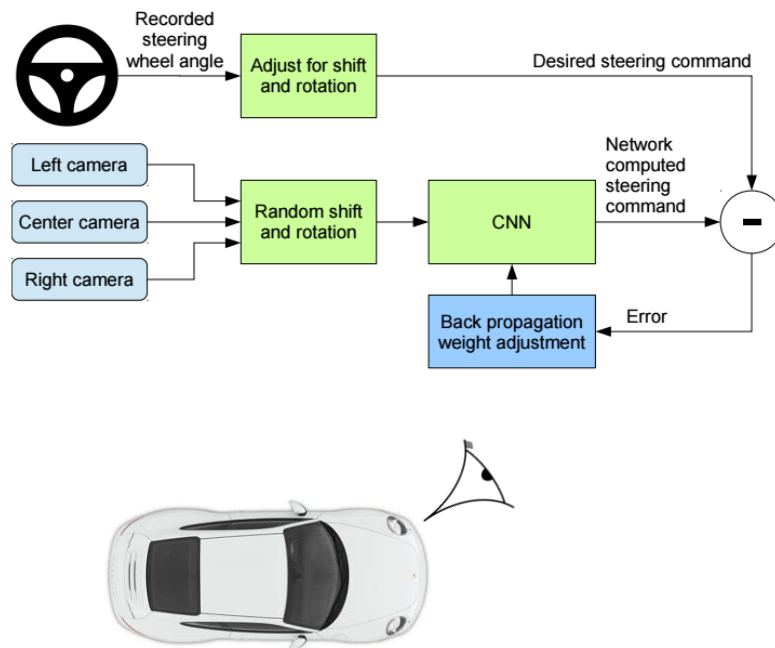
Why did that work?



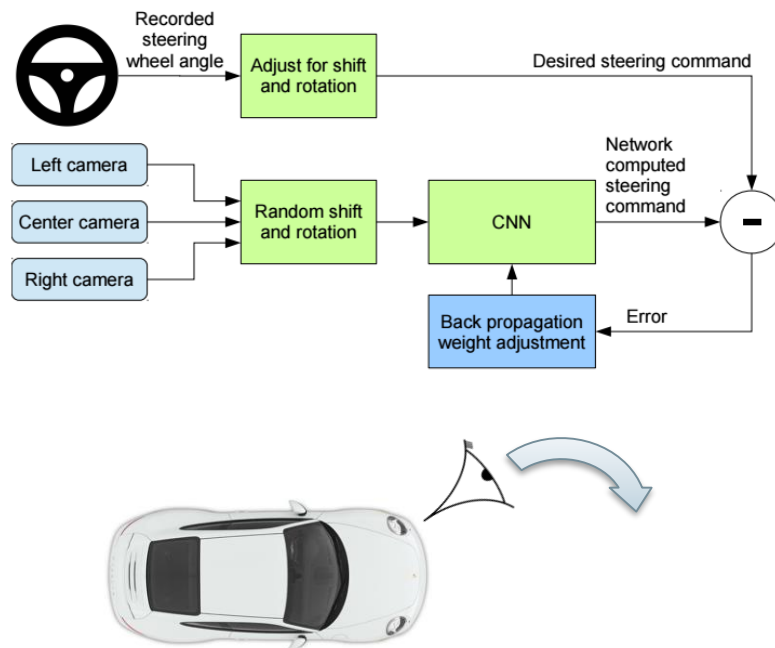
Why did that work?



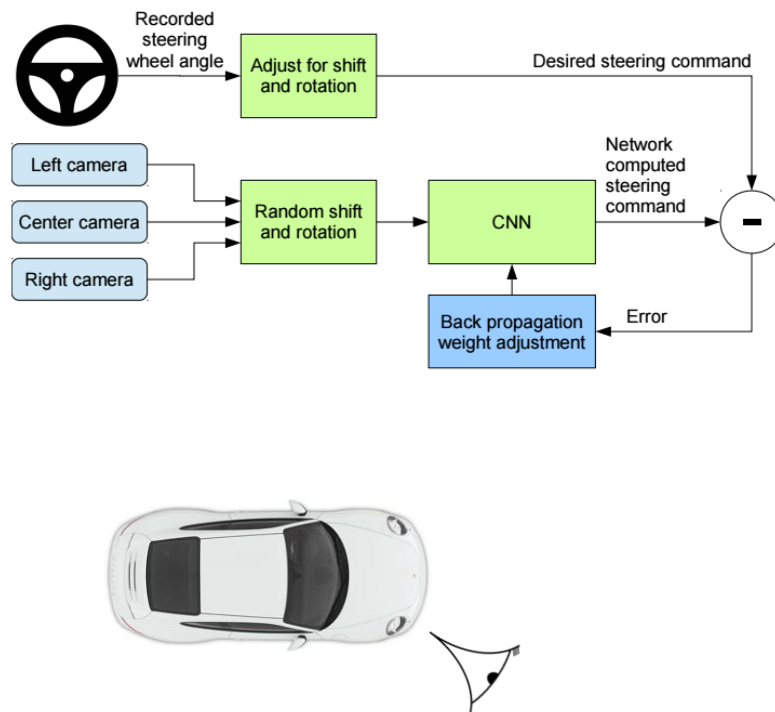
Why did that work?



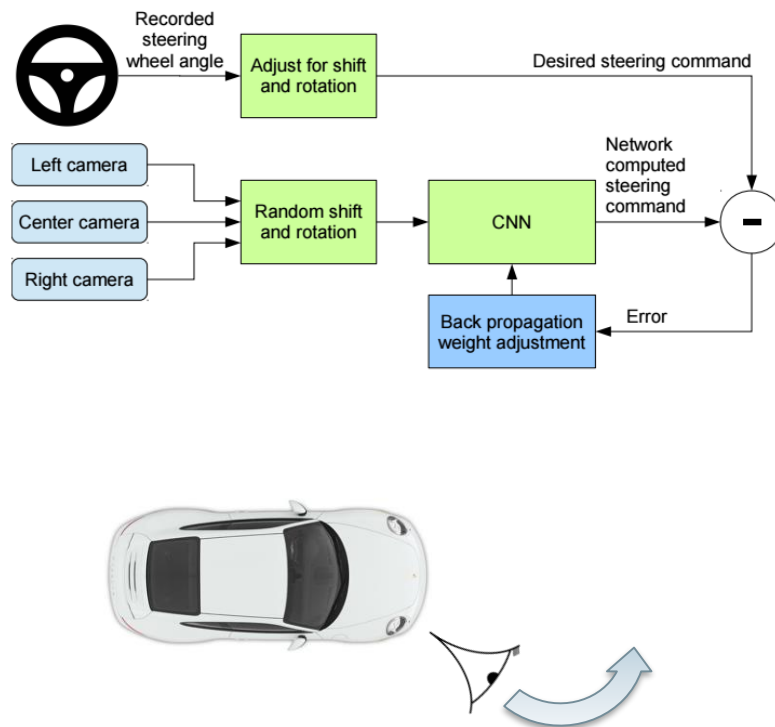
Why did that work?




Why did that work?




Why did that work?




Can we make it work more often?




Can we make it work more often?




Can we make it work more often?




Can we make it work more often?



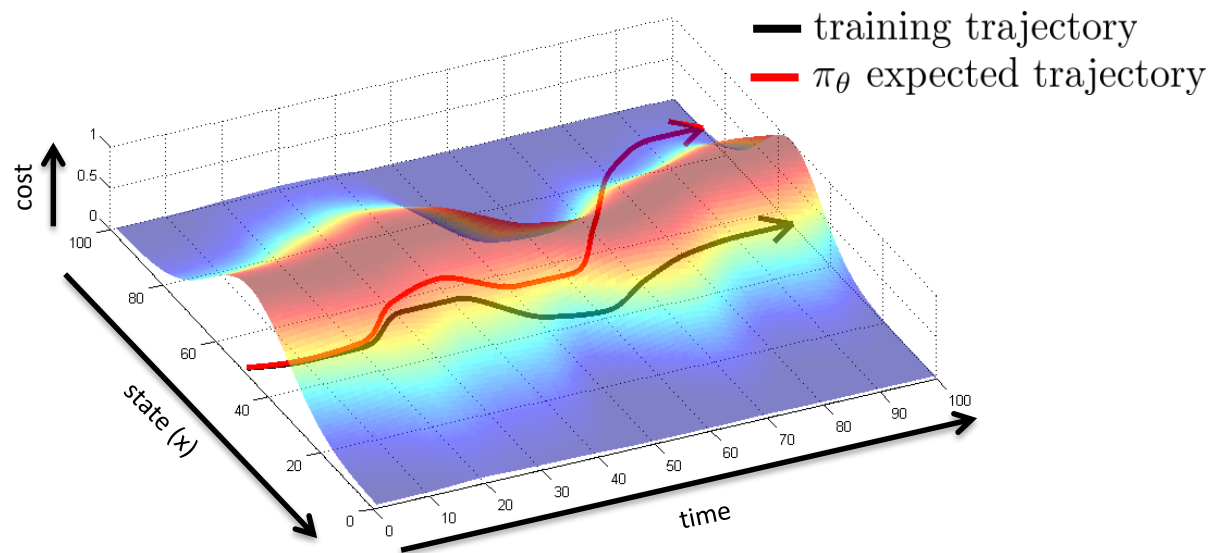
Can we make it work more often?




Can we make it work more often?




Can we make it work more often?




Can we make it work more often?



Can we make it work more often?




Can we make it work more often?




stability


Can we make it work more often?




Can we make it work more often?




Can we make it work more often?



Can we make it work more often?



Can we make it work more often?



can we make $p_{\text{data}}(\mathbf{o}_t) = p_{\pi_\theta}(\mathbf{o}_t)$?

More Terminology

Behavior Policy: The policy $\pi_{\theta}(u|o)$ that the agent uses to act in the world.



Target Policy: A policy $\pi_{\theta^t}^t(u|o)$ the agent is learning.




More Terminology

On Policy: Agent learns from its own experience, so target policy = behavior policy.



Off Policy: Target policy \neq behavior policy. More general. Can use experience from other agents



Can we make it work more often?

can we make $p_{\text{data}}(\mathbf{o}_t) = p_{\pi_\theta}(\mathbf{o}_t)$?

Current policy

$$\pi_\theta(u_t | o_t)$$



Can we make it work more often?

can we make $p_{\text{data}}(\mathbf{o}_t) = p_{\pi_\theta}(\mathbf{o}_t)$?

idea: instead of being clever about $p_{\pi_\theta}(\mathbf{o}_t)$, be clever about $p_{\text{data}}(\mathbf{o}_t)$!

Current policy

$$\pi_\theta(u_t | o_t)$$



Can we make it work more often?

can we make $p_{\text{data}}(\mathbf{o}_t) = p_{\pi_\theta}(\mathbf{o}_t)$?

idea: instead of being clever about $p_{\pi_\theta}(\mathbf{o}_t)$, be clever about $p_{\text{data}}(\mathbf{o}_t)$!

DAgger: Dataset Aggregation

Current policy

$$\pi_\theta(u_t | o_t)$$



Can we make it work more often?

can we make $p_{\text{data}}(\mathbf{o}_t) = p_{\pi_\theta}(\mathbf{o}_t)$?

idea: instead of being clever about $p_{\pi_\theta}(\mathbf{o}_t)$, be clever about $p_{\text{data}}(\mathbf{o}_t)$!

DAgger: Dataset Aggregation

goal: collect training data from $p_{\pi_\theta}(\mathbf{o}_t)$ instead of $p_{\text{data}}(\mathbf{o}_t)$

Current policy

$$\pi_\theta(u_t | o_t)$$



Can we make it work more often?

can we make $p_{\text{data}}(\mathbf{o}_t) = p_{\pi_\theta}(\mathbf{o}_t)$?

idea: instead of being clever about $p_{\pi_\theta}(\mathbf{o}_t)$, be clever about $p_{\text{data}}(\mathbf{o}_t)$!

DAgger: Dataset Aggregation

goal: collect training data from $p_{\pi_\theta}(\mathbf{o}_t)$ instead of $p_{\text{data}}(\mathbf{o}_t)$

how? just run $\pi_\theta(\mathbf{u}_t|\mathbf{o}_t)$

Current policy

$$\pi_\theta(u_t|o_t)$$



Can we make it work more often?

can we make $p_{\text{data}}(\mathbf{o}_t) = p_{\pi_\theta}(\mathbf{o}_t)$?

idea: instead of being clever about $p_{\pi_\theta}(\mathbf{o}_t)$, be clever about $p_{\text{data}}(\mathbf{o}_t)$!

DAgger: Dataset Aggregation

goal: collect training data from $p_{\pi_\theta}(\mathbf{o}_t)$ instead of $p_{\text{data}}(\mathbf{o}_t)$

how? just run $\pi_\theta(\mathbf{u}_t|\mathbf{o}_t)$

but need labels \mathbf{u}_t !

Current policy

$$\pi_\theta(u_t|o_t)$$



Can we make it work more often?

can we make $p_{\text{data}}(\mathbf{o}_t) = p_{\pi_\theta}(\mathbf{o}_t)$?

idea: instead of being clever about $p_{\pi_\theta}(\mathbf{o}_t)$, be clever about $p_{\text{data}}(\mathbf{o}_t)$!

DAgger: Dataset Aggregation

goal: collect training data from $p_{\pi_\theta}(\mathbf{o}_t)$ instead of $p_{\text{data}}(\mathbf{o}_t)$

how? just run $\pi_\theta(\mathbf{u}_t|\mathbf{o}_t)$

but need labels \mathbf{u}_t !

1. train $\pi_\theta(\mathbf{u}_t|\mathbf{o}_t)$ from human data $\mathcal{D} = \{\mathbf{o}_1, \mathbf{u}_1, \dots, \mathbf{o}_N, \mathbf{u}_N\}$

Current policy

$$\pi_\theta(u_t|o_t)$$



Can we make it work more often?

can we make $p_{\text{data}}(\mathbf{o}_t) = p_{\pi_\theta}(\mathbf{o}_t)$?

idea: instead of being clever about $p_{\pi_\theta}(\mathbf{o}_t)$, be clever about $p_{\text{data}}(\mathbf{o}_t)$!

DAgger: Dataset Aggregation

goal: collect training data from $p_{\pi_\theta}(\mathbf{o}_t)$ instead of $p_{\text{data}}(\mathbf{o}_t)$

how? just run $\pi_\theta(\mathbf{u}_t|\mathbf{o}_t)$

but need labels \mathbf{u}_t !

1. train $\pi_\theta(\mathbf{u}_t|\mathbf{o}_t)$ from human data $\mathcal{D} = \{\mathbf{o}_1, \mathbf{u}_1, \dots, \mathbf{o}_N, \mathbf{u}_N\}$
2. run $\pi_\theta(\mathbf{u}_t|\mathbf{o}_t)$ to get dataset $\mathcal{D}_\pi = \{\mathbf{o}_1, \dots, \mathbf{o}_M\}$

Current policy

$$\pi_\theta(u_t|o_t)$$



Can we make it work more often?

can we make $p_{\text{data}}(\mathbf{o}_t) = p_{\pi_\theta}(\mathbf{o}_t)$?

idea: instead of being clever about $p_{\pi_\theta}(\mathbf{o}_t)$, be clever about $p_{\text{data}}(\mathbf{o}_t)$!

DAgger: Dataset Aggregation

goal: collect training data from $p_{\pi_\theta}(\mathbf{o}_t)$ instead of $p_{\text{data}}(\mathbf{o}_t)$

how? just run $\pi_\theta(\mathbf{u}_t|\mathbf{o}_t)$

but need labels \mathbf{u}_t !

1. train $\pi_\theta(\mathbf{u}_t|\mathbf{o}_t)$ from human data $\mathcal{D} = \{\mathbf{o}_1, \mathbf{u}_1, \dots, \mathbf{o}_N, \mathbf{u}_N\}$
2. run $\pi_\theta(\mathbf{u}_t|\mathbf{o}_t)$ to get dataset $\mathcal{D}_\pi = \{\mathbf{o}_1, \dots, \mathbf{o}_M\}$
3. Ask human to label \mathcal{D}_π with actions \mathbf{u}_t

Current policy

$$\pi_\theta(u_t|o_t)$$



Can we make it work more often?

can we make $p_{\text{data}}(\mathbf{o}_t) = p_{\pi_\theta}(\mathbf{o}_t)$?

idea: instead of being clever about $p_{\pi_\theta}(\mathbf{o}_t)$, be clever about $p_{\text{data}}(\mathbf{o}_t)$!

DAgger: Dataset Aggregation

goal: collect training data from $p_{\pi_\theta}(\mathbf{o}_t)$ instead of $p_{\text{data}}(\mathbf{o}_t)$

how? just run $\pi_\theta(\mathbf{u}_t|\mathbf{o}_t)$

but need labels \mathbf{u}_t !

1. train $\pi_\theta(\mathbf{u}_t|\mathbf{o}_t)$ from human data $\mathcal{D} = \{\mathbf{o}_1, \mathbf{u}_1, \dots, \mathbf{o}_N, \mathbf{u}_N\}$
2. run $\pi_\theta(\mathbf{u}_t|\mathbf{o}_t)$ to get dataset $\mathcal{D}_\pi = \{\mathbf{o}_1, \dots, \mathbf{o}_M\}$
3. Ask human to label \mathcal{D}_π with actions \mathbf{u}_t
4. Aggregate: $\mathcal{D} \leftarrow \mathcal{D} \cup \mathcal{D}_\pi$

Current policy

$$\pi_\theta(u_t|o_t)$$



Can we make it work more often?

can we make $p_{\text{data}}(\mathbf{o}_t) = p_{\pi_\theta}(\mathbf{o}_t)$?


idea: instead of being clever about $p_{\pi_\theta}(\mathbf{o}_t)$, be clever about $p_{\text{data}}(\mathbf{o}_t)$!

DAgger: Dataset Aggregation

goal: collect training data from $p_{\pi_\theta}(\mathbf{o}_t)$ instead of $p_{\text{data}}(\mathbf{o}_t)$

how? just run $\pi_\theta(\mathbf{u}_t|\mathbf{o}_t)$

but need labels \mathbf{u}_t !

- 
1. train $\pi_\theta(\mathbf{u}_t|\mathbf{o}_t)$ from human data $\mathcal{D} = \{\mathbf{o}_1, \mathbf{u}_1, \dots, \mathbf{o}_N, \mathbf{u}_N\}$
 2. run $\pi_\theta(\mathbf{u}_t|\mathbf{o}_t)$ to get dataset $\mathcal{D}_\pi = \{\mathbf{o}_1, \dots, \mathbf{o}_M\}$
 3. Ask human to label \mathcal{D}_π with actions \mathbf{u}_t
 4. Aggregate: $\mathcal{D} \leftarrow \mathcal{D} \cup \mathcal{D}_\pi$

Current policy


$$\pi_\theta(u_t|o_t)$$




Dagger Example



What's the problem?


- 
1. train $\pi_{\theta}(\mathbf{u}_t|\mathbf{o}_t)$ from human data $\mathcal{D} = \{\mathbf{o}_1, \mathbf{u}_1, \dots, \mathbf{o}_N, \mathbf{u}_N\}$
 2. run $\pi_{\theta}(\mathbf{u}_t|\mathbf{o}_t)$ to get dataset $\mathcal{D}_{\pi} = \{\mathbf{o}_1, \dots, \mathbf{o}_M\}$
 3. Ask human to label \mathcal{D}_{π} with actions \mathbf{u}_t
 4. Aggregate: $\mathcal{D} \leftarrow \mathcal{D} \cup \mathcal{D}_{\pi}$

What's the problem?

- 
1. train $\pi_{\theta}(\mathbf{u}_t|\mathbf{o}_t)$ from human data $\mathcal{D} = \{\mathbf{o}_1, \mathbf{u}_1, \dots, \mathbf{o}_N, \mathbf{u}_N\}$
 2. run $\pi_{\theta}(\mathbf{u}_t|\mathbf{o}_t)$ to get dataset $\mathcal{D}_{\pi} = \{\mathbf{o}_1, \dots, \mathbf{o}_M\}$
 3. Ask human to label \mathcal{D}_{π} with actions \mathbf{u}_t
 4. Aggregate: $\mathcal{D} \leftarrow \mathcal{D} \cup \mathcal{D}_{\pi}$

What's the problem?


1. train $\pi_{\theta}(\mathbf{u}_t|\mathbf{o}_t)$ from human data $\mathcal{D} = \{\mathbf{o}_1, \mathbf{u}_1, \dots, \mathbf{o}_N, \mathbf{u}_N\}$
2. run $\pi_{\theta}(\mathbf{u}_t|\mathbf{o}_t)$ to get dataset $\mathcal{D}_{\pi} = \{\mathbf{o}_1, \dots, \mathbf{o}_M\}$
3. Ask human to label \mathcal{D}_{π} with actions \mathbf{u}_t
4. Aggregate: $\mathcal{D} \leftarrow \mathcal{D} \cup \mathcal{D}_{\pi}$



Domain Adaptation: Learning Reactive Controls for an MAV

Challenge: It's often much easier to get human training data in an environment different from the target environment (e.g. in simulation).

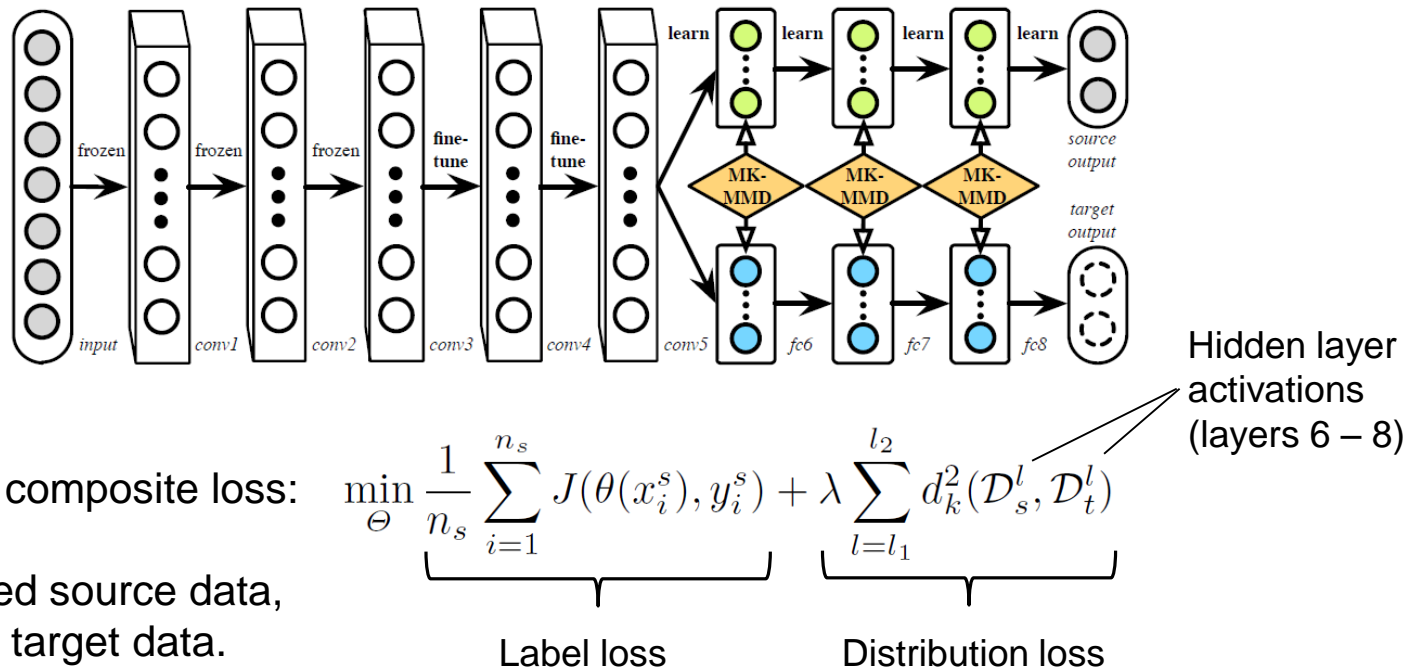
Developing a controller for the target domain after training in a different domain is a **domain adaptation** challenge.



Shreyansh Daftry, J. Andrew Bagnell, and Martial Hebert, “Learning Transferable Policies for Monocular Reactive MAV Control” 2016


Domain Adaptation: Learning Reactive Controls for an MAV

The domain adaptation network shares early layers, fine-tunes last CNN layers, and replicates FC layers:



Domain Adaptation: Learning Reactive Controls for an MAV

Examples:



Experiments and Results for (Row-1) Transfer across physical systems from ARDrone to ArduCopter, (Row-2) Transfer across weather conditions from summer to winter and (Row-3) Transfer across environments from Univ. of Zurich to CMU.


Reactive MAV Controls

Qualitative visualization of an example flight in dense forest.

The training data was collected from the same environment during summer season (Col-1) and tested during the winter season (Col-2).

The image sequence of MAVs on-board view is chronologically ordered from top to bottom and overlaid with color-coded commands issued by the policy learned using our proposed approach.

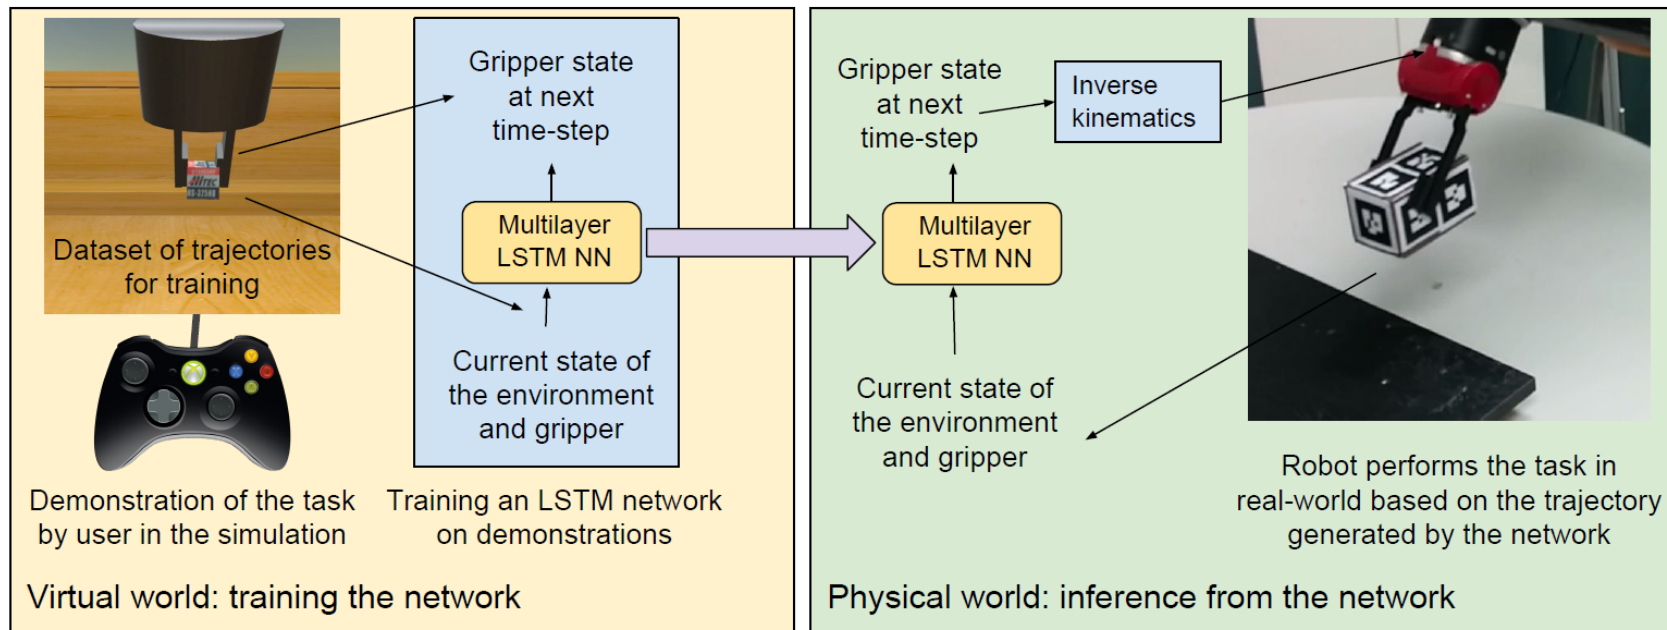
Additionally, we also compute the commands that would have been generated by the policy without domain adaptation (Col-3), for qualitative comparison.



Domain Adaptation: Learning Reactive Controls for an MAV

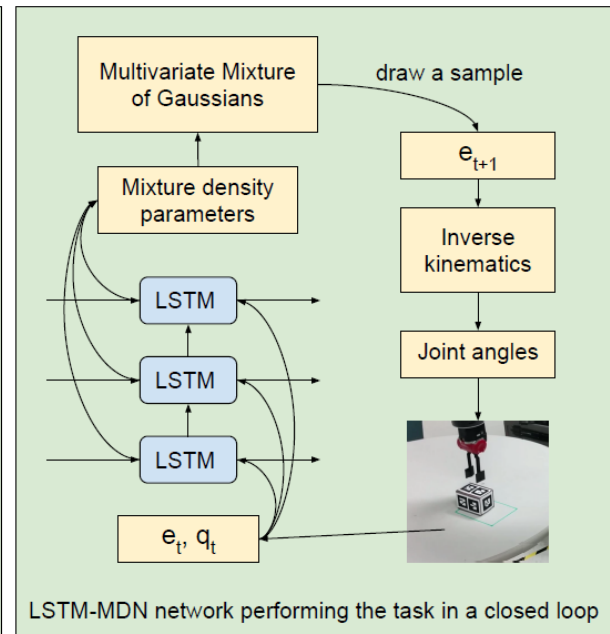
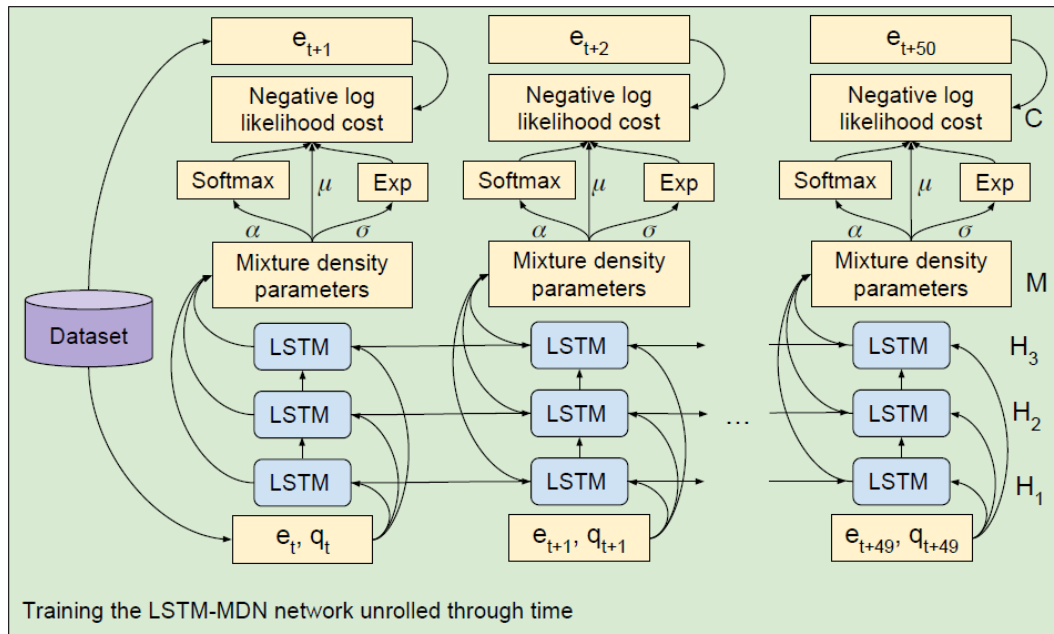
Video: <https://www.youtube.com/watch?v=Jvx0DWxTXAE>

Error Recovery: Learning ADL tasks with an LSTM



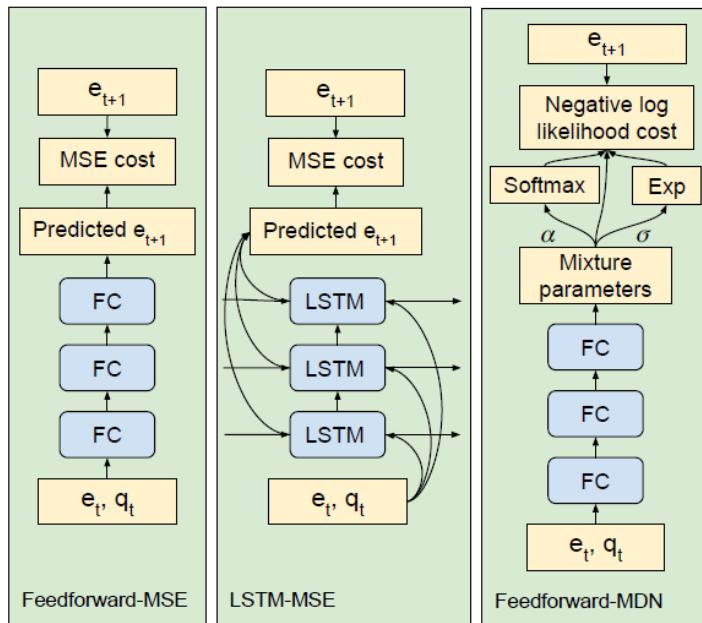
Rouhollah Rahmatizadeh, Pooya Abolghasemi, Aman Behal, Ladislau Boloni, "From Virtual Demonstration to Real-World Manipulation Using LSTM and MDN" 2016

Error Recovery: Learning ADL tasks with an LSTM



Error Recovery: Learning ADL tasks with an LSTM

Experiments: Comparison with baseline models (success rate):



Controller	Pick and place	Push to pose
Feedforward-MSE	0%	0%
LSTM-MSE	85%	0%
Feedforward-MDN	95%	15%
LSTM-MDN	100%	95%

Environment	Pick and place	Push to pose
Virtual world	100%	95%
Physical world	80%	60%


Error Recovery: Learning ADL tasks with an LSTM

Video: <https://youtu.be/9vYlIG2ozaM>

GAIL: Generative Adversarial Imitation Learning

Rather than trying to mimic the user blindly, try to **solve the same control problem that the user is solving**. i.e. estimate the user's **cost function**, and then optimize the cost by training.

This is called **Inverse Reinforcement Learning (IRL)**.



GAIL: Generative Adversarial Imitation Learning

Inverse Reinforcement Learning (IRL) is under-constrained, and often uses regularization heuristics.

Entropy regularization: define $H(\pi) \triangleq \mathbb{E}_{\pi}[-\log \pi(a|s)]$

Estimating the cost function is: $\max_{c \in \mathcal{C}} \left(\min_{\pi \in \Pi} -H(\pi) + \mathbb{E}_{\pi}[c(s, a)] \right) - \mathbb{E}_{\pi_E}[c(s, a)]$

Then the imitation learning problem is: $\text{RL}(c) = \arg \min_{\pi \in \Pi} -H(\pi) + \mathbb{E}_{\pi}[c(s, a)]$

GAIL: Generative Adversarial Imitation Learning

Inverse Reinforcement Learning (IRL) is under-constrained, and often uses regularization heuristics.

Entropy regularization: define $H(\pi) \triangleq \mathbb{E}_{\pi}[-\log \pi(a|s)]$

Entropy is highest for a random policy (random actions at every step).

Entropy is lowest (0) for a deterministic policy that takes a single action at each step.

GAIL: Generative Adversarial Imitation Learning

Inverse Reinforcement Learning (IRL) is under-constrained, and often uses regularization heuristics.

Entropy regularization: define $H(\pi) \triangleq \mathbb{E}_{\pi}[-\log \pi(a|s)]$

Estimating the cost function is: $\underset{c \in \mathcal{C}}{\text{maximize}} \left(\underset{\pi \in \Pi}{\min} -H(\pi) + \mathbb{E}_{\pi}[c(s, a)] \right) - \mathbb{E}_{\pi_E}[c(s, a)]$

Expert trajectory cost

Then the imitation learning problem is: $\text{RL}(c) = \arg \min_{\pi \in \Pi} -H(\pi) + \mathbb{E}_{\pi}[c(s, a)]$

GAIL: Generative Adversarial Imitation Learning

Inverse Reinforcement Learning (IRL) is under-constrained, and often uses regularization heuristics.

Entropy regularization: define $H(\pi) \triangleq \mathbb{E}_{\pi}[-\log \pi(a|s)]$

Estimating the cost function is: $\max_{c \in \mathcal{C}} \left(\min_{\pi \in \Pi} -H(\pi) + \mathbb{E}_{\pi}[c(s, a)] \right) - \mathbb{E}_{\pi_E}[c(s, a)]$

Learned policy cost

Then the imitation learning problem is: $\text{RL}(c) = \arg \min_{\pi \in \Pi} -H(\pi) + \mathbb{E}_{\pi}[c(s, a)]$

GAIL: Generative Adversarial Imitation Learning

MaxEnt IRL looks for a cost function which assigns low cost to the expert policy, and high cost to other policies.

Estimating the cost function is: $\underset{c \in \mathcal{C}}{\text{maximize}} \left(\min_{\pi \in \Pi} -H(\pi) + \mathbb{E}_{\pi}[c(s, a)] \right) - \mathbb{E}_{\pi_E}[c(s, a)]$

GAIL: Generative Adversarial Imitation Learning

GAIL then uses an adversary to discriminate the expert and learned policies by their state occupancy functions ρ_π and ρ_{π^E} . Don't worry about TRPO for now...

Algorithm 1 Generative adversarial imitation learning

- 1: **Input:** Expert trajectories $\tau_E \sim \pi_E$, initial policy and discriminator parameters θ_0, w_0
- 2: **for** $i = 0, 1, 2, \dots$ **do**
- 3: Sample trajectories $\tau_i \sim \pi_{\theta_i}$
- 4: Update the discriminator parameters from w_i to w_{i+1} with the gradient

$$\hat{\mathbb{E}}_{\tau_i} [\nabla_w \log(D_w(s, a))] + \hat{\mathbb{E}}_{\tau_E} [\nabla_w \log(1 - D_w(s, a))] \quad (17)$$

- 5: Take a policy step from θ_i to θ_{i+1} , using the TRPO rule with cost function $\log(D_{w_{i+1}}(s, a))$. Specifically, take a KL-constrained natural gradient step with



$$\begin{aligned} & \hat{\mathbb{E}}_{\tau_i} [\nabla_\theta \log \pi_\theta(a|s) Q(s, a)] - \lambda \nabla_\theta H(\pi_\theta), \\ & \text{where } Q(\bar{s}, \bar{a}) = \hat{\mathbb{E}}_{\tau_i} [\log(D_{w_{i+1}}(s, a)) \mid s_0 = \bar{s}, a_0 = \bar{a}] \end{aligned} \quad (18)$$

- 6: **end for**
-


GAIL: Generative Adversarial Imitation Learning

<https://www.youtube.com/watch?v=0hw0GD3IkA8>

Imitation learning: recap




Imitation learning: recap



Usually (but not always) insufficient by itself


Imitation learning: recap



Usually (but not always) insufficient by itself

Distribution mismatch problem

Imitation learning: recap




Usually (but not always) insufficient by itself

Distribution mismatch problem

Sometimes works well

Imitation learning: recap




Usually (but not always) insufficient by itself


Distribution mismatch problem

Sometimes works well

Micro-models (e.g. image and control transformations)



Imitation learning: recap




Usually (but not always) insufficient by itself

Distribution mismatch problem


Sometimes works well

Micro-models (e.g. image and control transformations)

Add more **on-policy** data, e.g. using Dagger



Imitation learning: recap



Usually (but not always) insufficient by itself

Distribution mismatch problem

Sometimes works well

Micro-models (e.g. image and control transformations)

Add more **on-policy** data, e.g. using Dagger

Domain adaptation and error recovery

